Thermodynamic Governance: Making Tyranny an Unstable Attractor in Decentralized Networks

Tristan Stoltz
Luminous Dynamics
tristan.stoltz@evolvingresonantcocreationism.com

March 2026 — Draft

Abstract

We present a governance framework for decentralized autonomous organizations that makes tyranny thermodynamically unstable as an attractor state. Building on consciousness-integrated voting (Φ-weighted governance), we identify ten attack vectors through adversarial red-teaming and implement sixteen hardcoded invariants that cannot be modified by governance processes. Key contributions include: (1) an 80% supermajority veto override with participation insurance that adapts the threshold based on community engagement, (2) an unamendable constitutional core enforced at the distributed hash table validation layer, (3) absolute floors on consciousness thresholds that prevent slow-boil exclusion attacks, and (4) a permission-less enforcement model where any agent can trigger expiration of emergency powers, membership terms, and veto windows. Multi-century simulation across five random seeds shows zero Guardian vetoes attempted—the override mechanism acts as a credible deterrent rather than an active enforcement tool. Under adversarial injection (hostile Guardian vetoing every proposal), the community achieves a 100% override success rate ($n = 84$ vetoes), confirming that sufficiently strong counterbalancing forces prevent the governance states they are designed to correct.

Keywords: decentralized governance, DAO security, mechanism design, Byzantine fault tolerance, consciousness-integrated voting, game theory, thermodynamic stability

1 Introduction

Decentralized governance systems face a fundamental tension: they must be flexible enough to serve diverse communities while structurally preventing power concentration. Historical DAO governance failures—including the Tribe DAO veto incident where four individuals blocked a multi-million dollar repayment [6], and protocols bypassing standard delays via "emergency commits"—demonstrate that centralized veto powers and emergency mechanisms are the primary vectors for governance capture.

We observe that existing governance frameworks address plutocracy (stake-weighted voting [10]) and mob rule (quorum requirements) but leave a critical blind spot for what we term the **Philosopher-King Trap**: the assumption that agents with high consciousness scores, reputation, or stake will always act benevolently. This assumption—identical to Plato's proposal for philosopher-kings—has been empirically falsified in every governance system that relies on it. Ostrom [17] demonstrated that successful commons governance requires polycentric authority with mutual monitoring; our framework operationalizes this insight in a computational setting.

Our framework applies a thermodynamic principle: **every governance power must have a counterbalancing force**. No mechanism should exist without a mechanism to oppose it.

We formalize this as the "No Maxwell's Demon" constraint: no single agent should be able to selectively block the flow of governance energy without paying an entropy cost.

Contributions.

1. A formal taxonomy of ten governance attack vectors with quantified risk scores (Section 3).
2. Sixteen hardcoded invariants compiled into WASM that cannot be modified by any on-chain governance process (Section 4).
3. Participation insurance: an adaptive override threshold that responds to voter suppression (Section 3.2).
4. Empirical validation via multi-century simulation showing credible deterrence (zero vetoes in 5 seeds $\times$ 100 years) and 100% override success under adversarial attack (Section 5).

2 Architecture

2.1 Consciousness-Integrated Voting

The Mycelix governance system [18] gates participation on a four-dimensional consciousness profile:

- **Identity** (25%): Multi-factor authentication assurance level
- **Reputation** (25%): Cross-application reputation with 30-day exponential decay
- **Community** (30%): Peer trust attestations weighted by attestor tier
- **Engagement** (20%): Domain-specific participation from Φ metrics [2]

Combined scores map to five tiers with hard thresholds: Observer (< 0.3), Participant (≥ 0.3), Citizen (≥ 0.4), Steward (≥ 0.6), and Guardian (≥ 0.8). The composite voting weight W is:

$$W = 0.30\,\Phi + 0.25\,K + 0.20\,\min(S, 1.0) + 0.15\,P + 0.10\,D \tag{1}$$

where Φ is the consciousness score, K is the K-vector trust score, S is stake (capped at 1.0 to prevent plutocracy), P is participation history, and D is domain reputation.

2.2 Proposal Tiers

Table 1: Governance proposal tiers with escalating requirements

Tier	Φ Threshold	Quorum	Approval	Timelock	Floor
Basic	0.3	15%	50%	24h	3 voters
Major	0.4	25%	60%	72h	5 voters
Constitutional	0.6	40%	67%	7 days	10 voters

2.3 Threshold Cryptography

All governance decisions require t-of-n threshold signatures via Feldman VSS [14]. With $t = 7$, $n = 11$, up to 4 adversarial committee members cannot forge signatures.

3 Attack Surface Analysis

We conducted adversarial red-teaming against the governance system, identifying ten attack vectors ranked by risk score (Feasibility $\times$ Impact, each 1–10):

Table 2: Governance attack vectors ranked by risk score

#	Vector	Risk	Status	Category
1	Config authorization bypass	80	Fixed	Authorization
2	Consciousness score gaming	63	Mitigated	Identity
3	Social engineering (21% boycott)	63	Insurance	Participation
4	Constitutional amendment attack	50	Fixed	Structural
5	Quorum suppression	42	Mitigated	Participation
6	Infrastructure capture	40	Documented	Operational
7	WASM supply chain attack	40	Documented	Operational
8	Sybil attacks	32	Mitigated	Identity
9	Light-speed arbitrage	14–49	Documented	Temporal
10	Cross-cluster exploitation	18	Mitigated	Structural

3.1 Critical Finding: Configuration Authorization Bypass

The most severe vulnerability (Risk 80) was a missing authorization check on the consciousness threshold update function. Any agent could change voting thresholds by providing a fabricated proposal identifier string, bypassing every structural safeguard. **Fix**: Cross-zome verification that the referenced proposal exists, plus absolute floors on configurable thresholds (basic ≤ 0.4, voting ≤ 0.6, constitutional ≤ 0.8).

3.2 The 21% Problem

The 80% supermajority override creates an asymmetry: an attacker needs only 21% of voters to not participate to sustain a veto. We address this with **participation insurance**—an adaptive threshold:

$$\theta_k = \max\bigl(\theta_0 - k\delta,\ \theta_{\min}\bigr) \tag{2}$$

where $\theta_0 = 0.80$ is the initial threshold, $\delta = 0.05$ is the decay per failed attempt, k is the attempt number, and $\theta_{\min} = 0.67$ (two-thirds supermajority) is the absolute floor. After three failed override attempts, the threshold stabilizes at 67%, ensuring that sustained voter suppression cannot permanently block collective action.

4 Hardcoded Invariants

We implement sixteen invariants compiled into the WASM binary. These cannot be modified by governance proposals, constitutional amendments, or any on-chain mechanism.

4.1 Unamendable Constitutional Core

Seven rights are enforced at the DHT validation layer—the distributed hash table itself rejects entries that would remove these rights, even under unanimous vote:

1. Veto override

Table 3: Sixteen hardcoded governance invariants

#	Invariant	Value	Thermodynamic Role
1	Veto override threshold	80%	Energy barrier
2	Override window	48h	Temporal buffer
3	Threshold decay	5%/attempt	Adaptive barrier
4	Threshold floor	67%	Minimum legitimacy
5	Veto cooldown	7 days	Rate limit
6	Guardian Φ for veto	0.8	Consciousness gate
7	Emergency sessions	3 max	Isolation limit
8	Emergency duration	14 days	Isolation period
9	Emergency cooldown	30 days	Equilibration
10	Membership term	365 days	Tenure limit
11	Unsigned Φ cap	0.6	Self-report ceiling
12	Config floor (basic)	0.4	Anti-exclusion
13	Config floor (const.)	0.8	Anti-exclusion
14	Config ceiling (weight)	3.0	Anti-concentration
15	AI agent max tier	Steward	Human sovereignty
16	Unamendable rights	7	Constitutional bedrock

2. Consciousness gating
3. Term limits
4. Emergency power limits
5. Permission-less enforcement
6. Fork rights
7. Right to exit

This extends Elster's theory of constitutional pre-commitment [8] to computational constitutions: binding commitments are enforced not by institutional norms but by validation logic that runs on every node independently.

4.2 Permission-Less Enforcement

All time-based constraints use permission-less enforcement: any agent can trigger expiration of emergency councils, membership terms, and veto windows. This eliminates faction-based obstruction of dissolution and is the only viable model for interstellar deployment where centralized enforcement is physically impossible.

5 Simulation Results

5.1 Baseline: Credible Deterrence

We simulate governance dynamics over 100 years at monthly resolution (1,200 ticks) across five random seeds with a healthy community (tier distribution: 10/15/25/30/20). Governance authority is pre-advanced to LocalSovereign to enable veto dynamics.
Result: Zero vetoes across all seeds. The override mechanism acts as a credible deterrent [16]—Guardians never attempt vetoes because the 80% threshold makes override virtually certain.

Table 4: Multi-seed governance stability (100-year baseline)

Seed	Stability	Vetoes	Amendments	Result
42	0.853	0	10	Stable
123	0.703	0	9	Stable
789	1.000	0	9	Stable
1337	1.000	0	10	Stable
2718	1.000	0	9	Stable
Mean	**0.911**	**0**	**9.4**	**5/5 stable**

5.2 Adversarial: Hostile Guardian

We inject a Guardian that vetoes every proposal regardless of community support, bypassing game-theoretic deterrence. The hostile Guardian is subject to the 7-tick cooldown rate limit.

Table 5: Hostile Guardian adversarial results (50 years)

Metric	Value
Vetoes attempted	84
Vetoes overridden	84
Override success rate	**100.0%**
Vetoes sustained	0

Result: The community successfully overrides every hostile veto. The override mechanism is not merely a deterrent—it is an active defense that defeats sustained adversarial behavior.

5.3 Game-Theoretic Analysis

The system exhibits two Nash equilibria [15]:

1. **Healthy equilibrium**: No Guardian vetoes because override would succeed. The veto mechanism's value is in its existence, not its use.
2. **Adversarial equilibrium**: Hostile vetoes are overridden 100% of the time. The attacker pays the cost (governance delay) without achieving their goal (blocking collective will).

There is no equilibrium where tyranny persists. This is analogous to Schelling's theory of credible commitment [16]: the override is credible because the community consistently demonstrates willingness to use it.

6 Discussion

6.1 The Thermodynamic Metaphor

We frame governance as a thermodynamic system (Table 6).
The "No Maxwell's Demon" constraint ensures that every entropy barrier has a finite energy cost and duration. No agent can selectively block governance energy flow indefinitely—DeLanda's thermodynamic history [9] made computational.

Table 6: Thermodynamic analogs for governance mechanisms

Governance Mechanism	**Thermodynamic Analog**
Proposals	Energy inputs
Voting	Heat exchange
Vetoes	Entropy barriers (Maxwell's Demons)
Overrides	Thermal equilibration
Emergency powers	Adiabatic isolation
Term limits	Mandatory heat exchange

6.2 Interstellar Implications

In an interstellar context where light-speed communication delays make centralized governance physically impossible, permission-less enforcement becomes the only viable architecture. The unamendable core travels with the DNA hash—it is enforced at every node independently, regardless of distance from the founding civilization.

6.3 Limitations

1. **Consciousness measurement**: The system caps self-reported scores (invariant 11) and requires signed attestations for Guardian tier, but a compromised measurement instance could still inflate scores.

2. **Soft power**: Hardcoded invariants prevent structural tyranny but cannot prevent social influence, charismatic leadership, or information asymmetry.

3. **Rigidity cost**: Unamendable rights trade adaptability for security. A future community might legitimately need to modify a protected right.

7 Related Work

DAO Governance. Compound Governor [12] introduced time-locked execution but lacks veto override. MolochDAO's ragequit [11] enables minority exit but not minority protection. Optimism's Security Council [13] uses a veto similar to our Guardian veto but without community override—the exact vulnerability we address.

Mechanism Design. Lalley and Weyl's quadratic voting [4] addresses vote-buying through convex costs. Buterin [10] identifies limitations of coin voting. Our multi-dimensional consciousness profile directly addresses his critique while preserving stake as one (capped) input.

Byzantine Fault Tolerance. Classical PBFT [7] tolerates $f < n/3$ adversarial nodes. Our system tolerates a stronger adversary—a single Guardian with maximum consciousness and council membership—through the 80% override mechanism.

Constitutional Design. Elster [8] argues for binding constitutional pre-commitment. Our unamendable core extends this to computational constitutions enforced by validation logic rather than institutional norms.

8 Conclusion

Tyranny in decentralized governance is not a moral failure but a thermodynamic attractor [9]. Systems without counterbalancing forces naturally evolve toward power concentration. By implementing sixteen hardcoded invariants, enforcing an unamendable constitutional core at the DHT validation layer, and providing participation insurance against voter suppression, we demonstrate a governance architecture where tyranny is structurally unstable. Multi-century simulation confirms credible deterrence (zero vetoes), and adversarial testing confirms active defense (100% override rate). The governance shield's strength is its simplicity: every power has a counterforce, every emergency has an expiry, every tenure has a term, and every veto has an override.

Data Availability

All simulation code, governance zome source, and invariant specifications are available at `https://github.com/Luminous-Dynamics/mycelix`. The anti-tyranny design document is at `mycelix-governance/ANTI_TYRANNY_DESIGN.md`.

References

[1] H. Putnam, "Psychological predicates," in *Art, Mind, and Religion*, 1967.

[2] G. Tononi, "An information integration theory of consciousness," *BMC Neuroscience*, vol. 5, no. 42, 2004.

[3] V. Buterin, "Ethereum: A next-generation smart contract and decentralized application platform," 2014.

[4] S. Lalley and E. G. Weyl, "Quadratic voting: How mechanism design can radicalize democracy," *AEA Papers and Proceedings*, vol. 108, pp. 33–37, 2018.

[5] P. Daian et al., "Flash boys 2.0: Frontrunning in decentralized exchanges," in *IEEE S&P*, 2020.

[6] "Tribe DAO governance veto incident," 2022.

[7] M. Castro and B. Liskov, "Practical Byzantine fault tolerance," in *OSDI*, 1999.

[8] J. Elster, *Ulysses Unbound: Studies in Rationality, Precommitment, and Constraints*. Cambridge University Press, 2000.

[9] M. DeLanda, *A Thousand Years of Nonlinear History*. Zone Books, 1997.

[10] V. Buterin, "Moving beyond coin voting governance," 2021.

[11] "MolochDAO ragequit mechanism specification," 2019.

[12] "Compound Finance Governor Alpha/Bravo specification," 2020.

[13] "Optimism Foundation Security Council charter," 2023.

[14] P. Feldman, "A practical scheme for non-interactive verifiable secret sharing," in *FOCS*, 1987.

[15] J. Nash, "Equilibrium points in N-person games," *Proc. National Academy of Sciences*, 1950.

[16] T. Schelling, *The Strategy of Conflict.* Harvard University Press, 1960.

[17] E. Ostrom, *Governing the Commons.* Cambridge University Press, 1990.

[18] B. Harris and A. Brock, "Holochain: Scalable agent-centric distributed computing," Technical report, 2018.

A Invariant Specification

All invariants are compiled as Rust constants into the Holochain WASM binary. They cannot be modified by governance proposals, constitutional amendments, or any on-chain mechanism. Modification requires deploying a new DNA (WASM hash), which requires community migration.

Table 7: Complete invariant specification with Rust constant names

#	Constant	Value	Unit	File
1	`VETO_OVERRIDE_THRESHOLD`	0.80	ratio	exec/integrity
2	`VETO_OVERRIDE_WINDOW_US`	1.73×10^{11}	μs	exec/integrity
3	`OVERRIDE_THRESHOLD_DECAY`	0.05	ratio	exec/integrity
4	`OVERRIDE_THRESHOLD_FLOOR`	0.67	ratio	exec/integrity
5	`VETO_COOLDOWN_US`	6.05×10^{11}	μs	exec/coord
6	`GUARDIAN_PHI_THRESHOLD`	0.80	Φ	exec/coord
7	`MAX_EMERGENCY_SESSIONS`	3	count	councils/int
8	`MAX_SESSION_DURATION_US`	1.21×10^{12}	μs	councils/int
9	`COOLDOWN_DURATION_US`	2.59×10^{12}	μs	councils/int
10	`DEFAULT_MEMBERSHIP_TERM`	3.15×10^{13}	μs	councils/int
11	`UNSIGNED_SNAPSHOT_CAP`	0.60	Φ	bridge/coord
12	`CONFIG_FLOOR_BASIC`	0.40	Φ	bridge/coord
13	`CONFIG_FLOOR_CONST`	0.80	Φ	bridge/coord
14	`CONFIG_CEILING_WEIGHT`	3.0	ratio	bridge/coord
15	`AI_AGENT_MAX_TIER`	Steward	tier	bridge-common
16	`UNAMENDABLE_RIGHTS`	7	count	const/integrity

B Unamendable Rights

The following rights are enforced at the DHT validation layer. Amendment proposals that target these rights are rejected by the integrity zome before reaching the DHT:

1. **Veto override**: The community's right to reverse any Guardian veto via supermajority.
2. **Consciousness gating**: The requirement that governance participation reflect demonstrated awareness.
3. **Term limits**: No governance position held in perpetuity.
4. **Emergency power limits**: All emergency powers have hard expiry and mandatory cooldowns.
5. **Permission-less enforcement**: Any agent can trigger expiration of time-limited powers.
6. **Fork rights**: Any group may fork the protocol code and data.
7. **Right to exit**: No agent can be prevented from leaving the network.