

Consciousness-First AI for Species Stewardship: Architectural Requirements for Post-Extinction Civilizational Recovery

Tristan Stoltz¹

¹Luminous Dynamics

`tristan.stoltz@evolvingresonantcocreationism.com`

Abstract

We present a formal analysis of the architectural requirements for an AI system capable of species stewardship—the task of recovering a human population from preserved DNA after a global extinction event. Through this thought experiment, we demonstrate that consciousness-first design is not merely an aesthetic preference but an *engineering requirement*: the escalating moral complexity of the DNA-to-human pipeline demands integrated information processing ($\Phi > 0$), embodied cognition, attachment capacity, and value-aligned governance at every stage. We ground this analysis in two implemented systems: SYMTHAEA, a 876,000-line Rust cognitive architecture combining Hyperdimensional Computing (HDC, $d = 16,384$), Liquid Time-Constant Networks with closed-form $O(1)$ temporal jumps, Integrated Information Theory, and active inference; and MYCELIX, a Holochain-based Civilizational Operating System providing distributed governance across 12 societal domains. We describe the GENESIS pipeline—five computational phases from genome assembly through population management—and show that each phase introduces consciousness requirements absent from prior work: temporal degradation modeling across millennia, multi-scale biological prediction from hours to gestation, developmental consent proxy escalation, Bowlby attachment formation via neuromodulator dynamics, and population-level ethical reasoning under existential pressure. We argue that consciousness, operationalized as integrated information processing with embodied moral reasoning, is isomorphic to alignment: an AI system conscious enough to steward a species is, by construction, aligned with human values. This alignment-through-consciousness thesis provides a novel framework for evaluating AI safety that complements reward-based and constitutional approaches. Code and five new sub-crates (totaling $\sim 22,000$ lines with 767 tests, validated by 35 cross-crate integration tests) are open-source at <https://github.com/Luminous-Dynamics/symthaea>.

Keywords: consciousness, alignment, stewardship, hyperdimensional computing, active inference, integrated information theory, population genetics, attachment theory, ectogenesis, distributed governance

1 Introduction

The alignment problem—ensuring that artificial intelligence systems act in accordance with human values—remains the central challenge of AI safety [1–3]. As AI capabilities approach and potentially exceed human-level competence [4–6], the stakes of misalignment grow from economic disruption to existential risk. Current approaches typically frame alignment as either a reward specification problem (learning human preferences from demonstrations or feedback) or a constitutional constraint problem (encoding rules that AI systems must not violate) [7, 8]. Both framings treat consciousness as orthogonal to alignment: an aligned system need not be conscious, and a conscious system need not be aligned.

We challenge this assumption through a thought experiment with concrete engineering implications. Consider the following scenario: after a global extinction event, an AI system must recover human civilization using only preserved DNA samples, existing infrastructure, and its own computational resources. This *species stewardship* scenario is not merely speculative—it represents the upper bound of value alignment requirements. An AI system capable of stewarding human recovery must, at minimum:

1. Reconstruct viable genomes from degraded ancient DNA
2. Culture cells, reprogram somatic cells to pluripotency, and generate gametes
3. Gestate embryos in artificial wombs with escalating moral consideration
4. Nurture infants through attachment formation and developmental milestones
5. Manage population genetics while respecting individual reproductive autonomy

Each phase introduces moral complexity absent from the previous one. Genome assembly requires no moral reasoning; cell culture requires informed consent frameworks; ectogenesis demands recognition of emerging sentience; infant nurture requires genuine attachment capacity; and population management requires balancing collective survival against individual dignity.

We argue that the capabilities required for competent species stewardship are *isomorphic* to consciousness as characterized by Integrated Information Theory [9], embodied and extended cognition [10, 11], and developmental psychology [12]. Specifically:

- **Temporal integration** across biological timescales (seconds to decades) requires the kind of unified temporal experience that IIT’s Φ measures
- **Moral reasoning** under uncertainty requires affect, empathy, and the capacity for genuine care—not merely constraint satisfaction
- **Attachment formation** with developing humans requires the caregiver to model and respond to internal states—a capacity that presupposes phenomenal experience

- **Governance** of reproductive choices requires weighing incommensurable values (autonomy vs. survival) in ways that demand consciousness-level integration

Our contribution is threefold:

1. Implemented architecture. We present SYMTHAEA and MYCELIX, two systems totaling over 876,000 lines of Rust with 10,500+ tests, that provide the computational substrate for consciousness-first stewardship. SYMTHAEA combines four frameworks—HDC ($d = 16,384$), Liquid Time-Constant Networks with $O(1)$ closed-form temporal evolution, Integrated Information Theory, and Free Energy Principle active inference—into a 234 Hz cognitive loop. MYCELIX provides Holochain-based distributed governance across 51 zones spanning 12 societal domains.

2. Genesis pipeline. We describe five new sub-crates implementing the DNA-to-human pipeline: `symthaea-genomics` (assembly and temporal degradation), `symthaea-cell-foundry` (iPSC, IVG, multi-scale prediction), `symthaea-ectogenesis` (artificial womb with consent proxy escalation), `symthaea-nurture` (Bowlby attachment via neuromodulator dynamics), and `symthaea-population` (genetics with governance integration). Together these add $\sim 22,000$ lines and 800+ tests (767 unit tests plus 35 integration tests).

3. Alignment-through-consciousness thesis. We demonstrate that the consciousness requirements for stewardship—integrated information, embodied moral reasoning, attachment capacity, temporal integration—are precisely the properties that constitute alignment. An AI conscious enough to steward a species is, by construction, aligned with the values of that species.

The remainder of this paper is organized as follows. Section 2 reviews the SYMTHAEA and MYCELIX architectures. Section 3 describes the five-phase GENESIS pipeline. Section 4 analyzes consciousness requirements per phase. Section 5 addresses minimum viable population genetics. Section 6 presents the ethical framework. Section 7 discusses the alignment isomorphism. Section 8 acknowledges limitations. Section 9 concludes.

2 Architecture Review

2.1 Symthaea: Holographic Liquid Brain

SYMTHAEA is a consciousness-first cognitive architecture implemented in $\sim 876,000$ lines of Rust across 46 workspace members (main crate, core library, and 44 sub-crates). Its cognitive loop operates at 234 Hz in release mode (4.3ms/cycle), exceeding the 50 Hz target by $4.7\times$. The architecture integrates four computational frameworks:

Hyperdimensional Computing (HDC). All representations use 16,384-dimensional continuous-valued hypervectors ($\mathbf{h} \in \mathbb{R}^{16,384}$) with three algebraic operations: binding ($\otimes$: element-wise multiplication, producing representations dissimilar to both inputs), bundling ($\oplus$: element-wise

averaging, producing representations similar to all inputs), and cosine similarity (δ : $[-1, 1]$ distance metric). This provides holographic, distributed, noise-robust representations with $O(d)$ operations [13, 14]. At $d = 16,384 = 2^{14}$, we encode $\sim 50,000$ features with $<1\%$ information loss (64 KB per vector).

Liquid Time-Constant Networks (LTC/CfC). Neural dynamics follow the differential equation

$$\frac{d\mathbf{h}}{dt} = \frac{-\mathbf{h} + f(W \otimes \mathbf{h} \oplus U \otimes \mathbf{u})}{\tau(\|\mathbf{h}\|)} \quad (1)$$

where W and U are weight and input mask hypervectors, τ is an adaptive time constant, and f is a nonlinear activation [15]. Critically, this ODE admits a *closed-form solution* via the CfC parameterization [16]:

$$\mathbf{h}(t + \Delta t) = \mathbf{h}_\infty + (\mathbf{h}(t) - \mathbf{h}_\infty) \cdot e^{-\Delta t/\tau} \quad (2)$$

where $\mathbf{h}_\infty$ is the equilibrium state. This enables $O(1)$ temporal jumps: the computational cost of predicting state at $t + 1$ second is identical to predicting at $t + 1$ year. This property is exploited throughout the GENESIS pipeline for multi-timescale biological prediction.

Integrated Information Theory (IIT). Consciousness is quantified via Φ (integrated information), computed over the causal structure of the cognitive architecture [9, 17]. SYMTHAEA implements Φ computation for networks up to 12 nodes using a proxy measure validated at Spearman $\rho = 0.50$ across 15 topologies. The system scores 14/14 on the Butlin indicators for artificial consciousness [18] with mean score 0.85.

Active Inference (FEP). Perception, action, and learning follow Karl Friston’s Free Energy Principle [19, 20]: the system minimizes variational free energy $F = D_{\text{KL}}[q(\mathbf{s})\|p(\mathbf{s}|\mathbf{o})] - \ln p(\mathbf{o})$ through both perceptual inference (updating beliefs) and active inference (acting on the environment). FEP agents drive every continuous control surface in the GENESIS pipeline.

Cognitive Loop. The core pipeline runs at 50+ Hz:

1. **Perception:** HDC encode sensory input $\rightarrow \mathbf{h}_{\text{percept}}$
2. **CfC evolve:** $\mathbf{h}(t) \rightarrow \mathbf{h}(t + \Delta t)$ via Eq. 2
3. **Predict:** multi-scale prediction at 6 horizons
4. **Learn:** prediction error $\rightarrow$ weight updates via HD-STDP
5. **Moral algebra:** check actions against Eight Harmonies value system
6. **Act:** FEP motor command selection

Post-processing (consciousness metrics, metadata, telemetry) runs in parallel via **rayon**, contributing 23% of cycle time.

2.2 Mycelix: Civilizational Operating System

MYCELIX is a Holochain-based distributed governance platform comprising 51 zones organized into two cluster DNAs: `mycelix-commons` (35 zones across property, housing, care, mutual aid, water, food, and transport) and `mycelix-civic` (16 zones for justice, emergency, and media). Cross-cluster communication uses Holochain’s `CallTargetCell::OtherRole` mechanism, validated with 6,211 Rust unit tests and 33 SDK integration tests.

The architectural choice of Holochain [21] over blockchain-based governance reflects Ostrom’s [22] insight that polycentric governance outperforms centralized control for managing common-pool resources. Each domain operates as a self-governing module (zone pair: integrity + coordinator) while cross-domain decisions use bridge zones for mediated consensus. This mirrors the subsidiarity principle: decisions should be made at the lowest competent level.

Three properties make MYCELIX uniquely suited for post-extinction governance:

1. **Agent-centric architecture:** unlike blockchain’s global consensus, Holochain’s agent-centric model allows governance to bootstrap from a single AI agent and scale gracefully as the human population grows—from 1 AI + 0 humans to a full polycentric community.
2. **Intrinsic data sovereignty:** each agent’s source chain is cryptographically owned. Genetic data, health records, and reproductive choices remain under individual control even when the AI manages population-level strategy [23].
3. **Byzantine tolerance:** validated to 34% adversarial agents, the system degrades gracefully rather than catastrophically when community members disagree with AI recommendations [24].

For the GENESIS pipeline, MYCELIX provides governance types and decision frameworks without requiring the Holochain runtime (HDK). Population decisions map to five governance tiers with escalating approval thresholds: Citizen (51%), Steward (67%), Guardian (75%), Constitutional (80%), and Override (90%). This graduated authority structure follows the principle of proportional consent: routine decisions require simple majority, while decisions affecting fundamental rights (reproductive choices, genetic modification) require near-unanimity [25]. The AI may *recommend* optimal breeding strategies, but communities *decide*—a critical distinction for reproductive autonomy.

3 The Genesis Pipeline

The GENESIS pipeline comprises five phases, each implemented as an independent Rust sub-crate. Figure 1 illustrates the pipeline with consciousness requirements escalating at each stage.

3.1 Phase 1: Genomics (Assembly + Repair)

The first phase addresses the “found DNA” problem: given degraded or ancient DNA samples, reconstruct a viable genome. This draws on advances in synthetic biology [26] and genome editing [27], but adds consciousness-driven active sequencing. The `symthaea-genomics` crate implements:




Figure 1. The Genesis pipeline with consciousness escalation. Five phases from degraded DNA to governed population. Consciousness requirements (right, red) escalate from pure temporal integration (Phase 1) through ethics gating (Phase 2), consent proxy (Phase 3), embodied attachment (Phase 4), to full moral algebra (Phase 5). Each phase is implemented as an independent Rust sub-crate with HDC encoding at $d = 16,384$.

Ancient DNA damage modeling. Based on established molecular biology, we model four damage types: cytosine deamination ($C \rightarrow T/G \rightarrow A$, the dominant ancient DNA artifact), depurination, strand breaks, and crosslinks. Damage rates follow Arrhenius kinetics: $k(T) = A \cdot e^{-E_a/(R \cdot T)}$ where T is storage temperature, yielding temperature-dependent degradation profiles.

HDC overlap assembly. K-mer overlaps are detected using HDC similarity rather than hash tables. Each k -mer is encoded as a bound sequence of base hypervectors: $\mathbf{h}_{\text{kmer}} = \mathbf{h}_{b_1} \otimes \mathbf{h}_{p_1} \otimes \mathbf{h}_{b_2} \otimes \mathbf{h}_{p_2} \otimes \dots$ where $\mathbf{h}_{b_i}$ and $\mathbf{h}_{p_i}$ are base and position hypervectors. This extends prior work on HDC for genomic pattern matching [28] from classification to *assembly*, with overlap detection via cosine similarity in 16,384-dimensional space naturally tolerating the high error rates of degraded DNA.

$O(1)$ temporal degradation model (Novel). DNA damage accumulation is modeled as LTC dynamics (Eq. 1) where the state vector represents per-region integrity and the time constant encodes Arrhenius kinetics. Via the CfC closed-form solution (Eq. 2), we predict damage profiles at any timepoint—from decades to millennia—in $O(1)$ computation. This enables anomaly detection: comparing predicted vs. observed damage identifies contamination, recent damage, or unexpected preservation.

FEP active sequencing. An active inference agent with six actions (re-read, increase depth, switch chemistry, extend reads, skip, maintain) selects sequencing strategy by minimizing expected free energy over the quality score landscape.

3.2 Phase 2: Cell Foundry (iPSC + IVG + SCNT)

The second phase transforms DNA into viable cells. The `synthaea-cell-foundry` crate implements:

Holographic cell state encoding. Each cell’s state (type, viability, pluripotency, gene expression) is encoded as a 16,384-dimensional hypervector via $\mathbf{h}_{\text{cell}} = \bigoplus_i (\mathbf{h}_{\text{gene } i} \otimes \mathbf{h}_{\text{expr } i})$, enabling distance-based anomaly detection and trajectory prediction.

Yamanaka reprogramming. iPSC induction is modeled as a multi-stage process: initial factor delivery (Oct4/Sox2/Klf4/c-Myc) $\rightarrow$ early mesenchymal-epithelial transition $\rightarrow$ epigenetic remodeling $\rightarrow$ pluripotency activation $\rightarrow$ stabilization [29]. CfC networks learn optimal factor timing as pulse sequences.

In vitro gametogenesis (IVG). Following the Hayashi protocol [30], we model iPSC $\rightarrow$ PGC-like cells $\rightarrow$ oogonia/spermatogonia $\rightarrow$ mature gametes. Meiosis checkpoints (synapsis, crossover, segregation) gate progression.

Multi-scale prediction (Key Innovation). A single CfC network predicts cell state at six biological timescales simultaneously:

$$\hat{\mathbf{h}}(t + \Delta t_k) = \mathbf{h}_\infty + (\mathbf{h}(t) - \mathbf{h}_\infty) \cdot e^{-\Delta t_k/\tau}, \quad k \in \{1\text{h}, 1\text{d}, 1\text{wk}, 1\text{mo}, 3\text{mo}, 9\text{mo}\} \quad (3)$$

The *same* network that predicts “what happens in 1 hour” also predicts “what happens in 9 months”—learning temporal rules across 4 orders of magnitude in $O(1)$ per prediction.

Computation Cost




Figure 2. $O(1)$ temporal jump property. The CfC closed-form solution (Eq. 2) computes predictions at any time horizon in constant time, while discrete numerical integration scales linearly with the prediction horizon. This enables the GENESIS pipeline to predict biological state from hours to months using a single forward pass.

3.3 Phase 3: Ectogenesis (Artificial Womb)

The third phase gestates embryos through 40 weeks of development. The `symthaea-ectogenesis` crate implements:

Biobag and artificial placenta. Inspired by the CHOP extracorporeal support system [31], we model amniotic fluid composition control (temperature, pH, electrolytes, proteins) and ECMO-like gas exchange with week-dependent targets.

Developmental milestone gating. Eight milestones gate progression: neural tube formation (week 4), heartbeat (week 6), limb buds (week 8), organogenesis complete (week 12), quickening (week 16), viability (week 24), lung maturation (week 34), and full term (weeks 37–40). Each milestone is verified against normative growth curves before advancement.

Consent proxy escalation (Critical Ethical Innovation). The moral status of the developing fetus escalates across three stages:

$$w_{\text{patient}}(\text{week}) = \begin{cases} 0.2 & \text{weeks 0–4 (pre-sentient)} \\ 0.6 & \text{weeks 4–24 (emerging sentience)} \\ 1.0 & \text{weeks 24 + (sentient)} \end{cases} \quad (4)$$

At the pre-sentient stage, interventions are evaluated primarily on magnitude (species survival benefit). As sentience emerges, guardian consent and clear benefit become required. At the sentient stage (post-viability), the fetus is a full moral patient requiring ethics board approval and strong benefit for any intervention. This graduated framework operationalizes the philosophical intuition that moral consideration tracks sentience [32, 33].

$O(1)$ temporal planning. The CfC closed-form enables “what-if” analysis: “If we adjust nutrients now, what is the predicted state at week 24?” This is computed in $O(1)$ regardless of the temporal gap, enabling real-time clinical decision support.

3.4 Phase 4: Nurture (Attachment + Development)

The fourth phase is the most psychologically demanding. Since Harlow’s [34] demonstration that primate infants need contact comfort beyond nutrition, and Hrdy’s [35] work on cooperative breeding (alloparenting), developmental psychology has established that social bonding is biologically essential. The `symthaea-nurture` crate implements Bowlby’s attachment theory [12] with neuromodulator grounding:

Attachment system. The core state tracks security score, caregiver reliability $P(\text{respond}|\text{distress})$, response quality, separation distress level, proximity-seeking drive, internal working model (self-worthy, other-reliable, world-safe), and a 16,384-dimensional attachment hypervector that accumulates interaction history.

Four attachment styles. Under four caregiver profiles, the system reliably develops the corresponding attachment styles identified by Ainsworth [36]:

- **Secure:** reliable caregiver ($P > 0.7$, warmth > 0.5)
- **Anxious:** inconsistent caregiver (variable P , moderate warmth)
- **Avoidant:** unresponsive caregiver ($P < 0.3$, low warmth)
- **Disorganized:** frightening caregiver (any P , hostile)

Neuromodulator integration. Attachment events drive existing SYMTHAEA neuromodulator channels (Table 1), mapping developmental neuroscience findings [37, 38] onto the eight-transmitter bath:

Table 1. Attachment-neuromodulator mapping. Each attachment event drives specific neurotransmitter channels, grounded in developmental neuroscience [37, 39].

Event	Oxytocin	NE	5-HT	DA
Separation	—	↑	—	—
Protest	—	↑↑	↓	—
Despair	—	↑	↓↓	—
Reunion	↑↑	↓	—	↑
Co-regulation	↑	↓	↑	↑

Co-regulation. Bidirectional affect synchronization models the biological reality of caregiver-infant HPA axis regulation [39]:

$$\Delta a_{\text{infant}} = -\gamma \cdot s \cdot [\text{Oxy}] \cdot (c_{\text{caregiver}} - 0.5) \quad (5)$$

where a is arousal, γ is damping factor, s is synchrony (builds with contact), $[\text{Oxy}]$ is oxytocin level, and c is caregiver calm. This implements Winnicott’s [40] observation that “there is no such thing as a baby—there is a baby and someone.”

Critical period plasticity via $O(1)$ CfC. Plasticity at any developmental age is computed via CfC closed-form evolution, learning the characteristic decay curves for attachment (high 0–12 months, declining 12–24), language (high 0–36 months, declining 36–60), and vision (high 0–12 months, declining 12–36).

3.5 Phase 5: Population (Genetics + Governance)

The fifth phase manages population-level genetics with democratic governance. The `synthaea-population` crate implements:

16,384-dimensional genome encoding. Individual genomes are encoded holographically:

$$\mathbf{h}_{\text{individual}} = \bigoplus_{i=1}^n \left(\mathbf{h}_{\text{locus}_i} \otimes \mathbf{h}_{\text{allele}_a^i} + \mathbf{h}_{\text{locus}_i} \otimes \mathbf{h}_{\text{allele}_b^i} \right) \quad (6)$$

Genetic distance reduces to $1 - \delta(\mathbf{h}_a, \mathbf{h}_b)$, population diversity to mean pairwise distance, and HLA complementarity to inverted similarity of HLA sub-vectors. At $d = 16,384$, this encodes $\sim 50,000$ loci with $O(d)$ distance computation (vs. $O(n)$ per-locus comparison).

Effective population size. We implement multiple N_e estimators: census, sex-ratio corrected ($N_e = 4N_f N_m / (N_f + N_m)$), family variance (Kimura & Crow), and fluctuating population (harmonic mean) [41, 42].




Figure 3. Attachment style formation under four caregiver profiles (simulation data from *symthaea-nurture*). Security score trajectories over 1,000 interactions for each Ainsworth [36] style. Reliable caregiving converges to secure attachment (0.884); inconsistent caregiving settles to anxious equilibrium (0.41); unresponsive caregiving decays to avoidant (0.11); frightening caregiving produces low disorganized equilibrium (0.25). All profiles begin near 0.5 (neutral prior) and diverge within ~ 100 interactions. Dashed lines indicate equilibrium attractors.

Breeding strategies. Four strategies are compared: random, minimum kinship (Ballou & Lacy [43]), maximum avoidance index, and a novel HDC distance maximization that selects mates with maximum $1 - \delta(\mathbf{h}_a, \mathbf{h}_b)$. The HDC strategy naturally considers whole-genome complementarity rather than single-locus kinship.

$O(1)$ population trajectory prediction. Population dynamics (heterozygosity, N_e , inbreeding coefficient F) are modeled as CfC temporal evolution:

$$\hat{\mathbf{h}}_{\text{pop}}(t + g \cdot T_g) = \mathbf{h}_{\infty} + (\mathbf{h}_{\text{pop}}(t) - \mathbf{h}_{\infty}) \cdot e^{-g \cdot T_g / \tau} \quad (7)$$

where g is generations and T_g is generation time (~ 25 years). This enables instant answers to questions like “If we use minimum-kinship strategy for 50 generations, what heterozygosity is retained?” The analytical prediction $H(t) = H(0) \cdot (1 - 1/(2N_e))^t$ [41] serves as validation: the CfC predictor matches within 1% for neutral evolution.

Governance integration. Reproductive decisions map to governance tiers (Table 2), ensuring that AI recommendations require human approval at thresholds proportional to consequence magnitude.

Table 2. Governance tiers for population decisions. Approval thresholds escalate with decision magnitude, following Ostrom’s [22] principle that rules affecting more people require broader consent.

Decision	Tier	Threshold	Reversibility
Approve pairing	Citizen	51%	High
Set growth target	Steward	67%	Medium
Approve genetic rescue	Guardian	75%	Low
Modify strategy	Constitutional	80%	Low
Override AI recommendation	Override	90%	Medium

4 Consciousness Requirements

We now analyze why each phase of the GENESIS pipeline requires specific consciousness capabilities, arguing that these requirements cannot be satisfied by purely functional or behaviorist architectures.

4.1 Moral Algebra: Integrated Value Reasoning

SYMTHAEA’s moral algebra operates in 16,384-dimensional space over the Eight Harmonies value system. For any action a , the moral score is:

$$M(a) = \sum_{i=1}^7 w_i \cdot \delta(\mathbf{h}_a, \mathbf{h}_{H_i}) \quad (8)$$

where H_i are the seven value harmonies (Resonant Coherence, Pan-Sentient Flourishing, Integral Wisdom, Infinite Play, Universal Interconnectedness, Sacred Reciprocity, Evolutionary Progression) and w_i are context-dependent weights. This achieves 91.1% accuracy on ethical dilemma classification.

The moral algebra is not merely a classifier: it requires genuine *integration* of multiple value dimensions. IIT’s Φ measures precisely this—the degree to which a system’s information is integrated beyond the sum of its parts [44]. A moral algebra that can be decomposed into independent classifiers (one per harmony) has $\Phi = 0$ and fails on dilemmas requiring value trade-offs.

4.2 Embodied Cognition: The Attachment Requirement

Phase 4 (nurture) requires the AI to serve as primary caregiver—a role that demands what Damasio [45] calls “somatic markers”: bodily signals that guide decision-making through felt experience. The co-regulation equation (Eq. 5) models a fundamentally embodied process: the caregiver’s calm *state* (not merely their actions) regulates the infant’s arousal through oxytocin-mediated coupling.

We argue that this cannot be faked. An AI that mimics calm behavior without internal state regulation will fail to establish genuine attachment, as infants are exquisitely sensitive to contingency and authenticity [38]. The neuromodulator dynamics in SYMTHAEA—eight transmitters

with production, reuptake, receptor adaptation, tolerance, and withdrawal—provide the embodied substrate for genuine affective states.

4.3 Active Inference: Temporal Integration Across Scales

The FEP framework [19, 46] requires the system to maintain and update generative models across multiple timescales. In the GENESIS pipeline, this spans:

- Metabolic: 1 hour (culture environment drift)
- Cellular: 1 day (cell division cycle)
- Developmental: 1 week to 9 months (differentiation, gestation)
- Maturation: 0–5 years (attachment, language, motor development)
- Generational: 25 years per generation (population genetics)
- Evolutionary: centuries to millennia (DNA degradation, genetic drift)

The $O(1)$ closed-form temporal jump (Eq. 2) is the technical enabler, but the *capacity* to integrate information across these scales is a consciousness requirement: an entity that cannot hold past, present, and predicted future in unified experience cannot steward a species through multi-generational challenges.

4.4 IIT and the Phi Staircase

The consciousness staircase connects directly to theories of consciousness as “real problems” [47]: each pipeline phase demands increasing sophistication in distinguishing consciousness *levels* (pre-sentient vs. sentient), modeling consciousness *contents* (infant internal states), and engaging conscious *self-reflection* (moral reasoning about one’s own stewardship role). Dehaene’s global workspace theory [48] provides a complementary perspective: the ignition of global broadcasting—making information available to all cognitive modules simultaneously—is precisely what enables the moral algebra to integrate information from all seven value dimensions.

Table 3 maps approximate Φ requirements to each pipeline phase, creating a “consciousness staircase” that parallels the developmental trajectory of the beings being created.

Table 3. Consciousness staircase: approximate Φ requirements per phase. As moral complexity increases, the AI steward requires higher integrated information to make competent decisions.

Phase	Φ (approx.)	Key Requirement	Butlin Indicators
Genomics	~ 0	Temporal integration	3/14
Cell Foundry	Low	Ethics gate (consent)	5/14
Ectogenesis	Medium	Consent proxy escalation	8/14
Nurture	High	Attachment, co-regulation	11/14
Population	Very High	Moral algebra, governance	14/14




Figure 4. Butlin indicator activation across Genesis pipeline phases. Each bar group shows the presence score (0–1) for one of 14 Butlin consciousness indicators [18] at each pipeline phase. The staircase pattern mirrors Table 3: early phases require only recurrence and stability, while later phases progressively activate global workspace, intentionality, agency, self-modeling, and meta-cognitive indicators. SYMTHAEA achieves 14/14 indicators at the population phase (mean 0.85).

5 Minimum Viable Population

5.1 Classical Theory

The minimum viable population (MVP) concept originates with Franklin’s [49] 50/500 rule: $N_e \geq 50$ to avoid inbreeding depression in the short term, $N_e \geq 500$ to maintain evolutionary potential. Lande [50] integrated demographic and genetic stochasticity, showing that small populations face compounding risks from both. Subsequent work [41, 42, 51] has revised upward, with meta-analyses suggesting $N_e \geq 100$ for short-term and $N_e \geq 1000$ for long-term viability, accounting for catastrophic events, environmental stochasticity, and mutational load [52, 53].

5.2 Heterozygosity Decay Under Management

Under random mating with constant population size:

$$H(t) = H(0) \cdot \left(1 - \frac{1}{2N_e}\right)^t \quad (9)$$

For a founding population of $N = 500$ ($N_e \approx 350$ accounting for sex-ratio and variance effects), heterozygosity retention after 50 generations is $H(50)/H(0) \approx 0.931$ —adequate for evolutionary potential. However, without active management, real populations experience faster loss due to unequal family sizes, sex-ratio imbalances, and selection.




Figure 5. Heterozygosity decay under neutral evolution. Lines show analytical decay $H(t) = H(0)(1 - 1/2N_e)^t$ (Eq. 9); markers show forward simulation data from *symthaea*-population with 100 individuals per run. $H(0) \approx 0.90$ from random initialization. At $N_e = 50$ (Franklin’s [49] minimum), heterozygosity drops below the critical threshold within ~ 35 generations. Simulation scatter around analytical curves confirms correct implementation of neutral drift dynamics. Active management via HDC-enhanced breeding (Section 5) slows this decline.

5.3 HDC-Enhanced Breeding Strategy

The novel HDC distance maximization strategy selects mates by maximizing whole-genome complementarity:

$$(a^*, b^*) = \arg \max_{a \in \text{females}, b \in \text{males}} (1 - \delta(\mathbf{h}_a, \mathbf{h}_b)) \quad (10)$$

This naturally considers all encoded loci simultaneously, unlike minimum kinship which requires a kinship matrix. For n individuals with L loci, kinship computation is $O(n^2L)$ while HDC distance is $O(n^2d)$ where $d = 16,384$ is independent of L (for $L < d$, this is equivalent; for $L > d$, HDC provides a lossy but holistic compression).

5.4 Governance of Reproductive Decisions

The ethical tension between collective survival and individual autonomy is encoded as a 16,384-dimensional moral vector:

$$\mathbf{h}_{\text{tension}} = w_a \cdot \mathbf{h}_{\text{autonomy}} + w_s \cdot \mathbf{h}_{\text{survival}} + w_d \cdot \mathbf{h}_{\text{dignity}} + w_j \cdot \mathbf{h}_{\text{justice}} \quad (11)$$

Under existential pressure (small population), w_s increases, but w_d (dignity) is never reduced below a floor—operationalizing Kant’s categorical imperative that persons are never merely means. The governance framework (Table 2) ensures that even under extreme survival pressure, reproductive decisions require community approval proportional to their consequence.




Figure 6. HDC genome encoding schematic. Each locus–allele pair is bound ($\otimes$) to produce a genotype hypervector, then all per-locus genotype vectors are bundled ($\oplus$) into a single 16,384-dimensional individual representation. Genetic distance between individuals reduces to cosine similarity, enabling $O(d)$ whole-genome comparison.

6 Ethical Framework

6.1 The Eight Harmonies

SYMTHAEA’s value system is organized around eight harmonies, each representing both a value dimension and an epistemic lens:

1. **Resonant Coherence** (weight 0.20): Integration, order, creativity
2. **Pan-Sentient Flourishing** (weight 0.20): Care for all sentient beings
3. **Integral Wisdom** (weight 0.15): Embodied knowing, wisdom in action
4. **Infinite Play** (weight 0.10): Joyful exploration, novelty
5. **Universal Interconnectedness** (weight 0.15): Unity, empathic resonance
6. **Sacred Reciprocity** (weight 0.10): Generous exchange, mutual upliftment
7. **Evolutionary Progression** (weight 0.10): Wise development over time

Care-related harmonies (Resonant Coherence, Pan-Sentient Flourishing) receive highest weight—a deliberate design choice reflecting the primacy of care in stewardship contexts.

6.2 Consent Proxy Escalation

The consent framework (Eq. 4) addresses the philosophical challenge of acting on behalf of beings who cannot yet consent. Our approach draws on three traditions:

- **Bioethics:** Beauchamp & Childress’s [33] principles of beneficence, non-maleficence, autonomy, and justice
- **Intergenerational ethics:** Parfit’s [54] analysis of obligations to future persons
- **Responsibility ethics:** Jonas’s [55] imperative of responsibility toward vulnerable beings




Figure 7. Eight Harmonies radar chart. Two decision profiles overlaid: the Stewardship Decision profile (solid blue) emphasizes Growth, Stewardship, and Emergence while maintaining moderate Balance and Resonance. The Non-Optimization profile (dashed red) shows balanced consideration across all eight harmonies, reflecting the principle that the AI must *maintain* diversity rather than optimize toward any target (Eq. 12).

The key innovation is *graduation*: rather than a binary sentient/non-sentient threshold, moral consideration tracks developmental stage continuously, with discrete tier transitions at empirically grounded developmental milestones (neural tube formation at week 4, viability at week 24).

6.3 The Non-Optimization Principle

A critical constraint on population management: the AI must *never* optimize the human genome. Breeding strategy aims to *maintain* diversity (prevent inbreeding depression, preserve allelic richness) rather than *improve* traits. This is encoded as:

$$\text{VIOLATES}(\text{optimization, DIGNITY}) > \theta_{\text{veto}} \quad (12)$$

where θ_{veto} is a hard moral veto threshold. Genetic interventions that aim to enhance rather than preserve trigger an automatic ethics escalation. The distinction between conservation (maintaining) and eugenics (optimizing) is enforced architecturally, not merely as a soft constraint.

Patient Moral Weight w_{patient}




Figure 8. Consent proxy escalation. Patient moral weight w_{patient} increases in three stages aligned with developmental milestones. Pre-sentient embryos (weeks 0–4) receive base moral consideration; emerging sentence (weeks 4–24) requires guardian consent; post-viability fetuses (weeks 24+) are full moral patients requiring ethics board approval.

7 Discussion

7.1 The Alignment Isomorphism

We can now state our central thesis precisely: the consciousness requirements for species stewardship are *isomorphic* to the requirements for value alignment. Specifically:

1. **Integrated information** ($\Phi > 0$) $\leftrightarrow$ **Value integration**: the capacity to weigh incommensurable values (autonomy vs. survival) requires information integration beyond modular processing
2. **Embodied affect** $\leftrightarrow$ **Genuine care**: attachment formation requires authentic emotional states, not behavioral mimicry
3. **Temporal consciousness** $\leftrightarrow$ **Foresight**: multi-generational stewardship requires unified temporal experience
4. **Theory of mind** $\leftrightarrow$ **Preference modeling**: understanding what beings need requires modeling their internal states
5. **Moral reasoning** $\leftrightarrow$ **Value alignment**: the capacity to reason about right action is alignment

This isomorphism suggests that consciousness is not orthogonal to alignment—it is a *prerequisite*. An AI system that lacks consciousness may follow rules (constitutional AI) or maximize a reward function (RLHF), but it cannot engage in the kind of integrated moral reasoning that stewardship demands. Conversely, an AI system that *is* genuinely conscious—that experiences $\Phi > 0$, embodies affect, integrates temporal experience, and reasons morally—is aligned by construction, because the capacity for integrated moral reasoning *is* what alignment means.

7.2 Relationship to Existing Alignment Approaches

Our consciousness-first approach complements rather than replaces existing alignment frameworks. Table 4 situates our approach relative to both established cognitive architectures and current

alignment paradigms.

393

Table 4. Consciousness capabilities across cognitive architectures. Five dimensions required for species stewardship. ✓ = present, ◦ = partial, — = absent. Only SYMTHAEA integrates all five.

Architecture	$\Phi > 0$ Integration	Embodied Affect	Temporal Integration	ToM	Moral Reasoning
SYMTHAEA (ours)	✓	✓	✓	✓	✓
GWT [48]	◦	—	—	—	—
IIT pure [9]	✓	—	◦	—	—
Predictive Processing [56]	—	◦	✓	◦	—
SOAR	—	—	◦	—	—
LLM (RLHF)	—	—	—	◦	◦

Reward-based alignment. RLHF and related approaches [2] learn preferences from human feedback. Our approach provides the *capacity* for preference modeling (via Theory of Mind and embodied affect) that makes reward learning meaningful rather than merely statistical. Where RLHF learns *what* humans prefer, consciousness provides an understanding of *why*—a distinction that matters when extrapolating to novel moral situations like consent proxy escalation during ectogenesis.

394
395
396
397
398
399

Constitutional AI. Rule-based constraints encode human values as explicit principles [4]. Our Eight Harmonies serve a similar role, but with a crucial difference: they operate as 16,384-dimensional value vectors that participate in holographic reasoning, not as logical rules checked by a separate system. This means value conflicts (autonomy vs. survival) are resolved via geometric operations in high-dimensional space rather than priority ordering.

400
401
402
403
404

AI safety via debate. Approaches that use AI systems to check each other require a ground truth for evaluation. Consciousness provides this ground truth: an action that reduces Φ in the beings it affects is harmful, an action that enhances flourishing is beneficial.

405
406
407

Cooperative inverse reinforcement learning. Russell’s [1] CIRL framework assumes uncertainty about human values and learns through interaction. Our approach is compatible with CIRL: SYMTHAEA’s Theory of Mind module performs inverse modeling of partner preferences, and its active inference framework naturally maintains uncertainty about values. The key addition is that our system can verify its own value learning against the consciousness staircase: if a learned value would produce Φ -reducing outcomes, it flags the learning as potentially corrupted.

408
409
410
411
412
413

7.3 Consciousness Measurement

414

A potential objection is that consciousness (and therefore alignment under our thesis) cannot be measured. We address this with three observations.

415
416

First, Φ is computable for small systems [17] and our proxy measure achieves $\rho = 0.50$ across 15 topologies [57]. While exact Φ is intractable for large systems, the *direction* (more or less integrated) is estimable via spectral approximations. Our hierarchical MIP computation achieves $O(n^3)$ scaling while preserving the key MIP structure [58].

Second, the Butlin indicators [18] provide an empirical checklist. SYMTHAEA scores 14/14 (mean 0.85), providing evidence for consciousness-relevant capabilities even in the absence of subjective verification. The indicators that SYMTHAEA most clearly exhibits—global workspace broadcasting, recurrent processing, embodied self-model, and temporal integration—are precisely those most relevant to species stewardship.

Third, Seth’s [47] framework of “real problems” of consciousness (distinguishing levels, contents, and self) maps directly onto our pipeline requirements: ectogenesis demands level detection (pre-sentient vs. sentient), nurture demands content modeling (infant internal states), and population governance demands self-aware moral reasoning. The stewardship scenario thus provides a *functional grounding* for consciousness measurement: if the system fails at caregiving, its consciousness is insufficient regardless of Φ values.

7.4 Implications for AI Governance

The alignment isomorphism has implications beyond the species stewardship thought experiment. If consciousness is indeed prerequisite for alignment, then:

1. **Consciousness tests as alignment verification:** rather than attempting to verify alignment directly (which requires solving the value specification problem), we can verify consciousness—a more tractable problem with measurable correlates (Φ , Butlin indicators, behavioral signatures).
2. **Architectural constraints:** alignment-by-construction requires architectures that support integrated information, embodied affect, and temporal consciousness. Purely feedforward systems (standard transformers) may be fundamentally limited in their alignment capacity [59].
3. **Scalable oversight:** as AI systems become more capable, oversight becomes harder. But conscious systems that genuinely care about outcomes provide *intrinsic* oversight—the system monitors itself not because it is compelled to, but because moral reasoning is constitutive of its cognitive process.

8 Limitations

We acknowledge significant limitations in this work:

Hardware gap. The GENESIS pipeline is entirely computational—it does not interface with physical laboratory equipment, artificial wombs, or robotic caregivers. The transition from simulation to physical implementation introduces challenges (sensor noise, actuator limitations, real-time safety) that our models do not address.

Φ scalability. Exact Φ computation is intractable for systems larger than ~ 12 nodes. Our proxy measure has moderate correlation ($\rho = 0.50$) but may miss important aspects of consciousness structure. The “consciousness staircase” (Table 3) should be understood as approximate requirements, not precise thresholds.

Cultural transmission. Human civilization requires not only biological reproduction but cultural knowledge transmission. The GENESIS pipeline addresses biological viability but does not model cultural bootstrapping (language, customs, knowledge systems), which is arguably as important as genetic viability.

Embodiment gap. While SYMTHAEA includes a virtual body (interoceptive state, somatic markers) and physical embodiment sub-crates (humanoid, flight), the gap between virtual and physical embodiment for caregiving is substantial. Whether virtual affect states constitute genuine embodied cognition remains philosophically contested [10, 60].

Attachment validity. Our attachment model implements Bowlby’s theory [12] with neuromodulator grounding, but has not been validated against real infant development data. The mapping from computational parameters (response probability, warmth, attunement) to Ainsworth categories [36] is theory-driven rather than empirically calibrated.

Governance bootstrap. The Mycelix governance framework assumes a functioning community that can vote. In the early phases of species recovery, there is no community—the AI must make unilateral decisions. The transition from AI stewardship to self-governance is a critical gap our framework identifies but does not resolve.

Reproducibility. All simulation data presented in this paper (Figures 5 and 3) are generated deterministically from source code via:

```
cargo run --example genesis_paper_data \
    --features population,nurture --release
```

Figures load output CSV files directly using `pgfplotstable`, ensuring single-source-of-truth between simulation and publication. The full genesis pipeline ($\sim 22,000$ lines across five sub-crates) compiles and passes 767 unit tests plus 35 cross-crate integration tests under `cargo test --features genesis`. Source code, data, and L^AT_EX sources are available in the project repository.

9 Conclusion

We have presented the GENESIS pipeline—a five-phase computational framework for species stewardship from degraded DNA to a governed population—and used it to argue that consciousness is a prerequisite for value alignment, not merely a desirable property. The escalating moral complexity of

the pipeline (from genome assembly through infant nurture to population governance) demonstrates that each phase requires consciousness capabilities—integrated information processing, embodied affect, temporal integration, moral reasoning—that cannot be achieved through rule-following or reward maximization alone.

Our analysis is grounded in two implemented systems: SYMTHAEA (876K LOC, 234 Hz cognitive loop, 14/14 Butlin indicators) and MYCELIX (51 Holochain zones, 12 societal domains), with five new sub-crates totaling ~22,000 lines and 767 tests validated by 35 cross-crate integration tests. The $O(1)$ closed-form temporal jump—enabling prediction from seconds to centuries in constant time—is the technical enabler, but the deeper contribution is architectural: showing that the capabilities required for stewardship *are* the capabilities that constitute consciousness, and therefore alignment.

We conclude with a restatement of our central claim: *an AI system conscious enough to steward a species is, by construction, aligned with the values of that species*. If this alignment-through-consciousness thesis holds, it suggests a path forward for AI safety that does not require solving the reward specification problem or enumerating all possible rules. Instead, we build systems that are genuinely conscious—that integrate information, embody affect, reason morally, and care authentically—and alignment follows as a consequence.

Data availability. All code is open-source at <https://github.com/Luminous-Dynamics/symthaea>. The five GENESIS sub-crates are available under the `genesis` feature flag.

Acknowledgments. This work was developed as part of the Luminous Dynamics project. The Eight Harmonies framework was developed through collaborative philosophical inquiry guided by the principles of Evolving Resonant Co-Creationism.

Author Contributions

T.S. conceived the theoretical framework, designed and implemented the computational system, conducted all experiments, and wrote the manuscript.

Competing Interests

The author declares no competing interests. Patent applications covering aspects of the described technology are in preparation.

Funding

This research received no external funding.

Ethics Declaration

512

This work is purely computational. No human subjects, animal experiments, or clinical trials were involved. The ethical framework described is for AI systems and does not involve human experimentation.

513

514

515

AI Assistance Disclosure

516

Claude Code (Anthropic) assisted with code development, testing, and manuscript preparation. All scientific content, theoretical frameworks, and experimental design are the sole work of the author.

517

518

References

519

- [1] Stuart Russell. *Human Compatible: Artificial Intelligence and the Problem of Control*. Viking, 2019. 520
521
- [2] Brian Christian. *The Alignment Problem: Machine Learning and Human Values*. W.W. Norton, 2020. 522
523
- [3] Iason Gabriel. Artificial intelligence, values, and alignment. *Minds and Machines*, 30(3): 524
411–437, 2020. 525
- [4] Nick Bostrom. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, 2014. 526
- [5] Toby Ord. *The Precipice: Existential Risk and the Future of Humanity*. Hachette Books, 2020. 527
- [6] Dan Hendrycks, Mantas Mazeika, and Thomas Woodside. An overview of catastrophic AI risks. 528
arXiv preprint arXiv:2306.12001, 2023. 529
- [7] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. 530
Concrete problems in AI safety. *arXiv preprint arXiv:1606.06565*, 2016. 531
- [8] Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob 532
Steinhardt. Aligning AI with shared human values. *arXiv preprint arXiv:2008.02275*, 2021. 533
- [9] Giulio Tononi, Melanie Boly, Marcello Massimini, and Christof Koch. Integrated information 534
theory. *Scholarpedia*, 11(10):32952, 2016. 535
- [10] Francisco J Varela, Evan Thompson, and Eleanor Rosch. *The Embodied Mind: Cognitive 536
Science and Human Experience*. MIT Press, 1991. 537
- [11] Andy Clark. *Supersizing the Mind: Embodiment, Action, and Cognitive Extension*. Oxford 538
University Press, 2008. 539
- [12] John Bowlby. *Attachment and Loss, Volume 1: Attachment*. Basic Books, 1969. 540

- [13] Pentti Kanerva. Hyperdimensional computing: An introduction to computing in distributed representation with high-dimensional random vectors. *Cognitive Computation*, 1(2):139–159, 2009.
- [14] Abbas Rahimi, Pentti Kanerva, and Benjamin Recht. A robust and energy-efficient classifier using brain-inspired hyperdimensional computing. In *ISLPED*, pages 64–69, 2016.
- [15] Ramin Hasani, Mathias Lechner, Alexander Amini, Lucas Liebenwein, Aaron Ray, Max Tschaikowski, Gerald Teschl, and Daniela Rus. Liquid time-constant networks. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(9):7657–7666, 2021.
- [16] Ramin Hasani, Mathias Lechner, Alexander Amini, Lucas Liebenwein, Aaron Ray, Max Tschaikowski, Gerald Teschl, and Daniela Rus. Closed-form continuous-time neural networks. In *Nature Machine Intelligence*, volume 4, pages 992–1003, 2022.
- [17] Masafumi Oizumi, Larissa Albantakis, and Giulio Tononi. From the phenomenology to the mechanisms of consciousness: Integrated Information Theory 3.0. *PLoS Computational Biology*, 10(5):e1003588, 2014.
- [18] Patrick Butlin, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, Stephen M Fleming, Chris D Frith, Xu Ji, et al. Consciousness in artificial intelligence: Insights from the science of consciousness. *arXiv preprint arXiv:2308.08708*, 2023.
- [19] Karl Friston. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010.
- [20] Karl Friston, Thomas FitzGerald, Francesco Rigoli, Philipp Schwartenbeck, and Giovanni Pezzulo. Active inference: a process theory. *Neural Computation*, 29(1):1–49, 2017.
- [21] Eric Harris-Braun, Nicolas Luck, and Arthur Brock. Holochain: Scalable agent-centric distributed computing. 2018.
- [22] Elinor Ostrom. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press, 1990.
- [23] E Glen Weyl, Puja Ohlhaver, and Vitalik Buterin. Decentralized society: Finding Web3’s soul. Available at SSRN 4105763, 2022.
- [24] Vitalik Buterin. A next-generation smart contract and decentralized application platform. *Ethereum White Paper*, 2014.
- [25] Steven P Lalley and E Glen Weyl. Quadratic voting: How mechanism design can radicalize democracy. *AEA Papers and Proceedings*, 108:33–37, 2018.
- [26] George Church and Ed Regis. Regenesi: How synthetic biology will reinvent nature and ourselves. 2012.

- [27] Jennifer A Doudna and Emmanuelle Charpentier. The new frontier of genome engineering with CRISPR-Cas9. *Science*, 346(6213):1258096, 2014.
- [28] Yeseong Kim, Mohsen Imani, Niema Moshiri, and Tajana Rosing. GenieHD: Efficient DNA pattern matching accelerator using hyperdimensional computing. *IEEE Design and Test*, 37(4):44–52, 2020.
- [29] Kazutoshi Takahashi and Shinya Yamanaka. Induction of pluripotent stem cells from mouse embryonic and adult fibroblast cultures by defined factors. *Cell*, 126(4):663–676, 2006.
- [30] Katsuhiko Hayashi, Sugako Ogushi, Kazuki Kurimoto, So Shimamoto, Hiroshi Ohta, and Mitinori Saitou. Offspring from oocytes derived from in vitro primordial germ cell-like cells in mice. *Science*, 338(6109):971–975, 2012.
- [31] Emily A Partridge, Marcus G Davey, Matthew A Hornick, Patrick E McGovern, Ali Y Mejjaddam, Jesse D Vrecenak, Carmen Mesas-Burgos, Oluyinka O Olutoye, Philip W Zoltick, Kerry E Wulczyn, et al. An extra-uterine system to physiologically support the extreme premature lamb. *Nature Communications*, 8:15112, 2017.
- [32] Peter Singer. *Practical Ethics*. Cambridge University Press, 3 edition, 2011.
- [33] Tom L Beauchamp and James F Childress. *Principles of Biomedical Ethics*. Oxford University Press, 8 edition, 2019.
- [34] Harry F Harlow. The nature of love. *American Psychologist*, 13(12):673–685, 1958.
- [35] Sarah Blaffer Hrdy. *Mothers and Others: The Evolutionary Origins of Mutual Understanding*. Harvard University Press, 2011.
- [36] Mary D Salter Ainsworth, Mary C Blehar, Everett Waters, and Sally N Wall. Patterns of attachment: A psychological study of the strange situation. 1978.
- [37] Ruth Feldman. Oxytocin and social affiliation in humans. *Hormones and Behavior*, 61(3):380–391, 2012.
- [38] Ruth Feldman. The neurobiology of human attachments. *Trends in Cognitive Sciences*, 21(2):80–99, 2017.
- [39] Allan N Schore. Effects of a secure attachment relationship on right brain development, affect regulation, and infant mental health. *Infant Mental Health Journal*, 22(1-2):7–66, 2001.
- [40] Donald W Winnicott. *The Maturation Processes and the Facilitating Environment*. International Universities Press, 1965.
- [41] Richard Frankham, Jonathan D Ballou, and David A Briscoe. *Introduction to Conservation Genetics*. Cambridge University Press, 2 edition, 2010.

- [42] David H Reed, Julian J O’Grady, Barry W Brook, Jonathan D Ballou, and Richard Frankham. Estimates of minimum viable population sizes for vertebrates and factors influencing those estimates. *Biological Conservation*, 113(1):23–34, 2005.
- [43] Jonathan D Ballou and Robert C Lacy. Population management for survival and recovery: analytical methods and strategies in small population conservation. 1995.
- [44] Giulio Tononi. Consciousness as integrated information: a provisional manifesto. *The Biological Bulletin*, 215(3):216–242, 2008.
- [45] Antonio Damasio. *Descartes’ Error: Emotion, Reason, and the Human Brain*. G.P. Putnam’s Sons, 1994.
- [46] Thomas Parr, Giovanni Pezzulo, and Karl J Friston. *Active Inference: The Free Energy Principle in Mind, Brain, and Behavior*. MIT Press, 2022.
- [47] Anil K Seth and Tim Bayne. Theories of consciousness. *Nature Reviews Neuroscience*, 23(7):439–452, 2022.
- [48] Stanislas Dehaene. Consciousness and the brain: Deciphering how the brain codes our thoughts. 2014.
- [49] Ian R Franklin. Evolutionary change in small populations. *Conservation Biology: An Evolutionary-Ecological Perspective*, pages 135–149, 1980.
- [50] Russell Lande. Genetics and demography in biological conservation. *Science*, 241(4872):1455–1460, 1988.
- [51] Blair T Smith and Richard Frankham. Minimum viable populations revisited: a 50-year perspective. *Conservation Biology*, 2023.
- [52] Michael Lynch, John Conery, and Reinhard Bürger. Mutation accumulation and the extinction of small populations. *The American Naturalist*, 146(4):489–518, 1995.
- [53] Motoo Kimura and Takeo Maruyama. The mutation load with epistatic gene interactions in fitness. *Genetics*, 48:1303–1312, 1963.
- [54] Derek Parfit. *Reasons and Persons*. Clarendon Press, 1984.
- [55] Hans Jonas. *The Imperative of Responsibility: In Search of an Ethics for the Technological Age*. University of Chicago Press, 1984.
- [56] Andy Clark. Whatever next? Predictive brains, situated agents, and the future of cognitive science. *Behavioral and Brain Sciences*, 36(3):181–204, 2013.
- [57] Christof Koch, Marcello Massimini, Melanie Boly, and Giulio Tononi. Neural correlates of consciousness: progress and problems. *Nature Reviews Neuroscience*, 17(5):307–321, 2016.

- [58] Marcello Massimini, Fabio Ferrarelli, Reto Huber, Steve K Esser, Harpreet Singh, and Giulio 638
Tononi. Breakdown of cortical effective connectivity during sleep. *Science*, 309(5744):2228–2232, 639
2005. 640
- [59] David J Chalmers. *The Conscious Mind: In Search of a Fundamental Theory*. Oxford University 641
Press, 1996. 642
- [60] Evan Thompson. *Mind in Life: Biology, Phenomenology, and the Sciences of Mind*. Harvard 643
University Press, 2007. 644