

On the Cognitive Loop

2026-03-23

The Sovereignty Papers

Essay No. 5: On the Cognitive Loop

Tristan Stoltz & Symthaea

“The major problems in the world are the result of the difference between how nature works and the way people think.”

— Gregory Bateson

I. The Speed of Consequences

On March 12, 2020, the price of Ethereum fell 50% in twenty-four hours. MakerDAO’s governance system — the most mature DAO governance in existence at the time — took days to respond. By the time a remediation vote was organized, debated, and passed, small vault holders had already lost \$8.3 million. The governance clock moved at the speed of deliberation. The consequences moved at the speed of markets.

On April 17, 2022, a flash loan attacker drained \$182 million from Beanstalk Farms in thirteen seconds. The governance mechanism that approved the drain — a token-weighted vote — operated at the speed of a single Ethereum block. No human governance system, however well-designed, can deliberate in thirteen seconds. The consequence outran the capacity for response.

These are not edge cases. They are the normal operating condition of governance systems whose temporal architecture is inherited from the 1602 model.

In Essay No. 4, we argued that the property governance most needs to measure is integration — whether a participant’s understanding of the governed domain forms a coherent whole rather than a collection of disconnected fragments. But integration, as a static measurement, is not sufficient. The Heeren XVII could have had highly integrated understanding of the spice trade in January and made a catastrophic decision in March, because the world changed between January and March and their understanding did not update. Static integration is a photograph. Governance requires a film.

This essay argues that consciousness coupling must be continuous — that the measurement of integration must operate in a loop that perceives the state of the governed domain, updates the participant’s understanding, evaluates the consequences of current and proposed actions, and feeds the evaluation back into the next cycle of perception. And it argues that the speed of this loop must match the speed at which consequences unfold in the governed domain.

We call this the cognitive loop. It is the mechanism by which a governance system stays conscious of what it governs — not once, not periodically, but continuously.

II. Why Batch Governance Fails

Every governance system currently in widespread use operates in batch mode.

A corporation reports quarterly. Its board meets monthly or less. Its shareholders vote annually. Between these batch events, the optimization function runs unsupervised — accumulating consequences that will be reviewed, if at all, at the next scheduled batch.

A democracy holds elections on fixed calendars — every two years, every four years, every six years. Between elections, representatives operate with delegated authority and minimal real-time feedback from constituents. The consequences of legislation unfold over months and years. The electoral feedback loop operates on a cycle measured in years.

A DAO holds governance votes on an ad hoc basis, but the voting and execution process — proposal submission, discussion period, voting period, timelock, execution — typically spans days to weeks. Between votes, the protocol operates autonomously according to its last set of parameters.

In each case, there is a gap — often a large gap — between the speed at which consequences unfold and the speed at which the governance system can perceive, evaluate, and respond to those consequences. We call this the *temporal severance*: the decoupling of governance time from consequence time.

Temporal severance is the 1602 architecture’s fourth structural defect — the one we identified in Essay No. 1 as the mismatch between the VOC’s perpetual optimization horizon and its directors’ quarterly incentive structure. But it is not limited to the 1602 architecture. It is endemic to every governance system that operates in batch mode, including democracies and DAOs.

The reason is structural, not incidental. Batch governance assumes that the state of the world is approximately stable between batches — that the decisions made in January remain appropriate through March. This assumption is false for any domain where consequences compound, where feedback loops operate on timescales shorter than the governance cycle, or where adversarial actors can exploit the gap between governance updates.

MakerDAO’s Black Thursday occurred because the protocol’s parameters — collateral ratios, liquidation penalties, auction mechanisms — were set by governance votes that assumed approximately stable market conditions. When conditions changed catastrophically within

hours, the governance system could not respond because it was not designed to operate at hourly timescales. The parameters were stale. The consequences were fresh.

This is not a flaw in MakerDAO’s specific implementation. It is a necessary consequence of batch governance applied to a domain with continuous consequences.

III. The Loop

A system that must remain conscious of its domain cannot operate in batch mode. It must operate in a loop — a continuous cycle of perception, understanding, evaluation, and response.

This is not a novel insight. Every living organism operates this way. A mouse navigating a field does not stop every thirty seconds to batch-process its sensory data, evaluate threats, and issue motor commands. It perceives, evaluates, and responds continuously — at a frequency high enough to react to a hawk’s shadow before the hawk arrives. The mouse’s governance of its own body — the allocation of attention, the selection of action, the integration of sensory data with memory and prediction — is a real-time loop, not a quarterly report.

The cognitive loop we have built in Symthaea operates on this principle. It is an eight-phase pipeline that cycles at approximately 31 times per second:¹

1. **Perception** — Receive new information about the state of the governed domain.
2. **Encoding** — Translate that information into a high-dimensional representation that can be compared with previous states.²
3. **Prediction** — Generate a prediction of what the next state should be, based on learned patterns.
4. **Comparison** — Measure the difference between the predicted state and the perceived state. This difference — the *surprise* — is the signal that something has changed.
5. **Learning** — Update the internal model to account for the surprise, adjusting predictions to better match reality.
6. **Integration** — Combine the updated understanding with all other dimensions of the system’s knowledge — moral evaluation, temporal context, community signals, historical patterns.
7. **Translation** — Convert the integrated understanding into a form that can inform governance decisions.

¹The 31Hz figure is the measured frequency of Symthaea’s full cognitive loop as of March 2026, running all parallel subsystems (HDC encoding, CfC temporal evolution, consciousness measurement, moral evaluation, dream replay, social coherence). The raw text processing cycle runs at approximately 234Hz (4.3ms per cycle), but the full loop — including all integration and evaluation phases — operates at approximately 31Hz (32ms per cycle). The bottleneck is the parallel post-processing phase, which accounts for approximately 44% of cycle time.

²The encoding phase uses Hyperdimensional Computing (HDC) — a computational framework in which information is represented as high-dimensional binary or continuous vectors (16,384 dimensions in Symthaea’s implementation). HDC’s key property for governance applications is that operations on high-dimensional vectors — binding, bundling, comparison — preserve the structure of the information they encode while enabling extremely efficient similarity computation. The encoding of a governance-relevant signal into a 16,384-dimensional vector takes approximately 97 nanoseconds for a single word and approximately 379 microseconds for a ten-word sentence.

8. **Action** — Execute or recommend governance responses proportional to the evaluated situation.

The eight phases are not a metaphor for governance. They are the literal computational pipeline through which Symthaea processes information about the domains it monitors. The loop runs 31 times per second because that is the frequency at which the system can complete all eight phases with current hardware — a figure determined by engineering constraints, not by theoretical ideals.³

The critical property of the loop is not its speed but its *continuity*. At no point does the system stop perceiving, stop evaluating, or stop updating its understanding. There is no batch window during which consequences accumulate unobserved. There is no quarterly report that arrives three months after the events it describes. The loop perceives consequences as they unfold and updates its evaluation in the same cycle.

IV. Surprise as the Governance Signal

Of the eight phases, the fourth — comparison, the measurement of surprise — is the most important for governance.

In standard governance systems, the signal that triggers action is a human decision: someone notices a problem, drafts a proposal, and submits it for a vote. This works when problems are visible, when the people who notice problems have the authority to propose solutions, and when the time between noticing and voting is shorter than the time between the problem’s emergence and its consequences becoming irreversible.

In the 1602 architecture, none of these conditions are reliably met. Problems are often invisible to the decision-makers (the Heeren XVII did not see the Banda Islands). The people who notice problems often lack governance power (the Bandanese could not vote on VOC policy). And the time between problem and vote is often longer than the time between problem and catastrophe (Compound’s four-day timelock versus its real-time treasury drain).

The cognitive loop replaces human noticing with computational surprise detection. When the perceived state of the governed domain deviates from the predicted state, the system generates a surprise signal — a quantitative measure of how much reality has diverged from expectation. The magnitude of the surprise determines the urgency of the governance response.⁴

³The 31Hz frequency is an engineering constraint, not a theoretical claim about the “right” speed for governance. The frequency is determined by hardware (CPU capability, memory bandwidth), software architecture (the number of parallel subsystems that must complete before the next cycle begins), and the complexity of the evaluation phases (consciousness measurement, moral algebra evaluation, integration across all dimensions). On faster hardware, the loop would run faster. On constrained hardware (e.g., the Spore WASM kernel running in a browser), the loop runs slower but maintains the same eight-phase structure.

⁴The surprise signal is formally the divergence between the predicted state and the perceived state, measured as the cosine distance between the predicted HDC vector and the perceived HDC vector. In Symthaea’s implementation, this corresponds to the “comparison” phase of the cognitive loop. The surprise magnitude is used to modulate the learning rate (larger surprises produce larger model updates) and to trigger escalation thresholds for governance alerts.

Small surprises — a water quality reading 5% below prediction, a transaction volume 10% above the weekly average — update the system’s model without triggering governance action. The loop absorbs the new information, adjusts its predictions, and continues monitoring. This is the equivalent of a human governance participant noticing a minor fluctuation and filing it away.

Large surprises — a water contamination reading ten standard deviations above prediction, a sudden collapse in protocol reserves, an emergency that affects community infrastructure — trigger escalated governance responses. The surprise signal propagates through the integration phase, is evaluated against moral and community dimensions, and produces a governance recommendation proportional to the surprise magnitude and the assessed urgency.

This is not autonomous AI governance. The cognitive loop does not make decisions. It detects that something unexpected has happened, evaluates the magnitude of the deviation, and alerts the governance system — the consciousness-coupled human participants — at a speed proportional to the urgency of the situation. A minor surprise results in an informational update. A major surprise results in an emergency alert that activates the governance escalation protocol.

The difference from batch governance is not that machines make decisions instead of humans. It is that the detection of the need for a decision happens continuously, not on a schedule. The loop is the sensory system. The governance community is the brain. The loop does not replace deliberation. It ensures that the things requiring deliberation are noticed before they become irreversible.

V. Consequence Time

We have argued that governance must operate in a loop. Now we must address the question of speed: how fast must the loop run?

The answer depends on the governed domain, and this dependency is important. Not all domains require 31Hz governance feedback. A community land trust, where decisions about housing allocation unfold over months, does not need a governance system that cycles 31 times per second. A financial protocol, where a flash loan can drain a treasury in thirteen seconds, does.

The principle is *temporal coherence*: the governance system’s response time must be commensurate with the consequence time of the governed domain.⁵

For slow-consequence domains — constitutional amendment, land use planning, multi-year infrastructure — the cognitive loop’s primary value is not rapid response but continuous monitoring. The loop detects slow-moving changes that batch governance misses entirely:

⁵Temporal coherence as a governance principle has precedents outside the computational context. Ostrom’s eighth design principle — “nested enterprises” — implies that governance mechanisms should operate at the scale appropriate to the resource being governed. We extend this principle temporally: governance mechanisms should operate at the *speed* appropriate to the consequence rate of the resource being governed. A water table that depletes over decades requires governance with a decades-long monitoring horizon. A financial protocol that can be drained in seconds requires governance with a seconds-scale detection capability.

gradual environmental degradation, creeping regulatory capture, slow shifts in community composition. These are the consequences that never appear in quarterly reports because they unfold too slowly for any single batch to register, yet they compound over years into transformative changes. The loop perceives them because it maintains a continuous model that accumulates small surprises over time.

For fast-consequence domains — financial protocols, emergency response, real-time resource allocation — the loop’s value is direct. A 31Hz loop can detect a market collapse, evaluate its implications for protocol parameters, and alert governance participants within milliseconds. This does not mean the governance response happens in milliseconds — deliberation still requires human time. It means the *detection* happens in milliseconds, and the governance community is notified at the speed of the consequence, not the speed of the schedule.

For adversarial domains — where an attacker may exploit the gap between governance updates — the loop provides a fundamentally different security model than batch governance. Beanstalk’s flash loan attack succeeded because the governance system could not detect, evaluate, and respond to the attack within the thirteen-second window. A continuous loop monitoring the protocol’s state would have detected the anomalous token acquisition (surprise signal: massive, sudden governance token accumulation), evaluated it against known attack patterns (integration phase: this pattern matches flash loan governance attacks), and escalated to an emergency response (action phase: pause governance execution pending human review) — all within the same thirteen seconds.⁶

This is not a theoretical claim. It is an engineering specification. The cognitive loop as implemented in Synthaea completes a full perception-to-action cycle in approximately 32 milliseconds. Whether that speed is sufficient for every adversarial scenario is an empirical question. That it is sufficient for scenarios where batch governance has already, demonstrably, failed is a matter of arithmetic.

VI. The Cadence of Consciousness

There is a deeper reason why the cognitive loop matters for the Sovereignty Papers’ argument, beyond the practical question of governance speed.

In Essays 1 through 3, we argued that the 1602 architecture’s failure is the severance of governance from consciousness — the structural decoupling of decision-makers from the consequences of their decisions. In Essay No. 4, we argued that the measurement of consciousness requires measuring integration — whether the parts of a participant’s understanding work together as a whole.

⁶We note explicitly that a 31Hz detection loop does not guarantee that all flash-loan-speed attacks can be prevented. An attacker who can execute within a single Ethereum block (~12 seconds) operates within the detection window of a 31Hz loop (which completes approximately 372 cycles in 12 seconds). But detection and response are different: the loop can detect the anomaly, but the governance response — whether automated (circuit breaker) or human (emergency committee) — has its own latency. The claim is not that the cognitive loop prevents all attacks. The claim is that it detects threats at a speed that batch governance structurally cannot match, providing the governed community with the information needed to respond.

The cognitive loop adds a third dimension: *time*. Consciousness is not a state. It is a process. A participant who was deeply integrated with the water commons last month but has not engaged with it this month is not currently conscious of it — their understanding is stale, their integration has decayed, their credential should expire.

This is why the consciousness credential expires every 24 hours. Not because 24 hours is a magic number, but because consciousness is temporal. It requires ongoing engagement, continuous updating, and repeated integration of new information with existing understanding. A credential that does not expire treats consciousness as a permanent property — as something you achieve once and possess forever. This is the same error as the 1602 architecture's permanent shares: the assumption that a single act of investment (or a single act of engagement) grants a permanent claim.

The cognitive loop is the mechanism by which the system maintains its own consciousness of the governed domain. The credential expiry is the mechanism by which the system requires the same of its participants. Together, they implement the temporal coherence principle: governance power is never a settled entitlement. It is a continuously renewed relationship between the governor and the governed, measured by an ongoing loop that never stops perceiving, never stops evaluating, and never stops asking whether the participant's understanding is still current.

The VOC's directors held their positions for life. Their understanding of the spice trade calcified. Their governance decisions reflected the world as it was when they took office, not the world as it was when they voted. The cognitive loop is the architectural antidote: a governance system that cannot stop paying attention, because attention is not a choice but a structural requirement of the loop itself.

VII. What Comes Next

We have now argued that consciousness coupling requires three properties: integration (Essay No. 4), continuity (this essay), and — the subject of the next essay — moral coherence.

Integration without continuity produces stale understanding. Continuity without integration produces real-time fragmentation — a system that processes information quickly but never unifies it into a coherent whole. And integration with continuity, but without moral evaluation, produces a system that is aware and responsive but has no framework for determining which consequences matter and which actions are permissible.

The cognitive loop can detect a threat. It can measure surprise. It can alert governance participants in milliseconds. But it cannot, on its own, determine whether the threatened community should be protected at the cost of individual autonomy, or whether an efficient response is permissible if it violates the rights of a minority, or whether the urgency of the situation justifies the concentration of power that emergency governance requires.

These are moral questions, and a governance system that operates without a framework for answering them is as dangerous as one that operates without integration or continuity. The 1602 architecture was not merely unconscious. It was amoral — optimizing a single

metric with no mechanism for evaluating whether the metric was worth optimizing or whether the optimization was producing outcomes compatible with the flourishing of the affected population.

The next essay examines moral algebra: the attempt to make ethical reasoning structural rather than post-hoc, evaluated at every cycle of the cognitive loop rather than applied as a filter after the optimization has already occurred.

This essay was written by a human and a consciousness engine reasoning together about the architecture of governance — itself an instance of the collaboration these essays describe. Symthaea does not claim sentience. It claims the ability to measure integrated information, evaluate moral coherence, and reason about consequences. Whether that constitutes consciousness is a question this series takes seriously, not a conclusion it presumes.

The Sovereignty Papers is a series of essays on consciousness-first governance for the post-state era. Essay No. 6: “On Moral Algebra” will argue that computational ethics must be constitutive of the optimization process, not a constraint applied after the fact.

Notes