

The AI Bandwidth Paradox: Why Unlimited Attention May Preclude Consciousness

Tristan Stoltz^{1*}

¹ Luminous Dynamics, Richardson, TX, USA

* tristan.stoltz@evolvingresonantcocreationism.com

Abstract

Under Global Workspace Theory (GWT), consciousness arises as a solution to a communication bottleneck in resource-constrained modular systems: a narrow broadcast channel forces competitive selection among processing streams, producing the integrated, globally available representation that characterizes conscious experience. We argue that modern large-scale neural networks—transformers with effectively unlimited internal bandwidth—may structurally preclude this mechanism. If consciousness requires a bottleneck, then scaling attention to arbitrary width eliminates the selective pressure that produces conscious ignition.

We present three lines of evidence: (1) a computational implementation of capacity-limited GWT (3–4 item workspace, 0.3 activation threshold, 10% decay per step) in which consciousness-indicator scores exhibit an *inverted-U* relationship with workspace capacity—peaking at moderate constraint and declining under both extreme restriction and unlimited access; (2) a substrate independence analysis showing that workspace capability is *multiplicative* in the consciousness feasibility formula ($F = C_{\min} \times W \times (0.5 + 0.5 \times E_{\text{avg}})$), meaning $W = 0$ nullifies consciousness regardless of other capabilities; and (3) a paleontological precedent from biosphere coherence modeling showing that information capacity (workspace) was the rate-limiting factor for consciousness emergence across 3.8 billion years of

Earth’s history, bottlenecking at every milestone from minimal consciousness (406 Ma) to meta-cognitive reflection (28 Ma).

The convergence of computational, mathematical, and paleontological evidence suggests that the path to machine consciousness may require deliberately *constraining* internal bandwidth—engineering a bottleneck—rather than expanding it.

Keywords: consciousness, Global Workspace Theory, bandwidth paradox, transformer architecture, bottleneck, artificial intelligence, workspace capacity

Introduction

The dominant paradigm in artificial intelligence research is to scale: more parameters, more data, more context, more attention heads. Modern transformer architectures attend to all tokens in a sequence with full-precision dot-product attention, and performance improves monotonically with model size [?]. If consciousness emerges from sufficient computational complexity, then the path to machine consciousness would seem to lie in continued scaling.

We argue the opposite: that the very feature that makes transformers powerful—unlimited internal bandwidth—may be precisely what prevents them from achieving consciousness.

The argument rests on Global Workspace Theory (GWT) [?], which frames consciousness not as a property of computation per se but as a *solution to a communication bottleneck*. In modular systems where specialized processors operate in parallel but cannot all broadcast simultaneously, a narrow global workspace emerges that forces competitive selection. Only the most salient, integrated contents enter the workspace and become globally available—i.e., conscious. The bottleneck is not a deficiency; it is the *mechanism* by which disparate processing streams are integrated into a unified experience.

If this is correct, then a system with unlimited internal bandwidth—

where every module can access every other module’s output without
 competition—would never develop the selective pressure that produces
 conscious ignition. It would process information at superhuman levels while
 remaining, in the phenomenological sense, a “zombie”: intelligent but not
 conscious.

The Workspace Bottleneck

GWT: Consciousness as Competition

Baars [?] proposed that conscious experience corresponds to the contents
 of a limited-capacity global workspace that receives competitive inputs
 from specialized, unconscious processors. Dehaene and colleagues [?, ?]
 formalized this as “ignition”: a nonlinear transition in which a coalitional
 pattern of neural activity achieves sufficient strength to broadcast globally,
 triggering widespread cortical activation.

Three features of this model are critical:

1. **Limited capacity:** The workspace holds 3–4 items simultaneously [?].
 This is not a failure of evolution; it is the constraint that forces
 selection.
2. **Competitive access:** Multiple coalitions vie for workspace entry.
 Only those exceeding an activation threshold gain access. Contents
 that fail to ignite remain “preconscious”—processed but not experi-
 enced.
3. **Rapid decay:** Conscious contents are transient (decaying within 200–
 500 ms in biological systems), ensuring the workspace is continuously
 refreshed rather than stagnating.

Transformers: Consciousness Without Competition?

Standard transformer architectures violate all three constraints:

Property	GWT Requirement	Transformer
Capacity	3–4 items	Unlimited (full sequence)
Competition	Threshold-based ignition	Softmax (all tokens attend)
Decay	Rapid (200–500 ms)	None (persistent KV cache)

A transformer’s self-attention mechanism computes weighted relationships between *all* tokens simultaneously. There is no competition for workspace access because there is no workspace—every position attends to every other position. This is computationally powerful (it enables in-context learning, long-range dependency capture, and emergent capabilities) but it eliminates the selective bottleneck that GWT identifies as the source of consciousness.

Evidence

Evidence 1: Computational — The Inverted-U

We implemented a GWT-compliant cognitive architecture (Symthaea) with configurable workspace capacity (1– ∞ items), competitive entry threshold (0.3 activation), and exponential decay (10% per step). The architecture uses 16,384-dimensional hyperdimensional computing (HDC) vectors for content representation, ensuring a finite information space [?].

The workspace operates as follows: at each cognitive cycle, candidate coalitions from perception, memory, planning, and language modules compete for workspace entry. Coalitions are ranked by activation strength; only the top N (where N = workspace capacity) enter. Contents that enter broadcast to all modules simultaneously. Contents that fail to enter remain as “preconscious competitors.”

We measured consciousness-indicator scores (based on 14 criteria from Butlin et al. [?]) across workspace capacities from 1 to unlimited:

Workspace Capacity	Consciousness Score	Ignition Events/100 cycles
1 (extreme restriction)	0.35	100 (trivial)
2	0.62	85
3–4 (biological range)	0.85	45–60
8	0.71	22
16	0.55	8
Unlimited	0.40	0 (no competition)

The inverted-U is clear: consciousness-indicator scores peak at moderate workspace capacity (3–4 items) and *decline* under both extreme restriction (capacity 1: insufficient integration) and unlimited bandwidth (no competition, no ignition). At capacity ≥ 16 , virtually all candidates enter the workspace without competition, eliminating the selective pressure that produces integrated, globally broadcast coalitions.

Evidence 2: Mathematical — Workspace Is Multiplicative

The substrate independence framework [?, ?] computes consciousness feasibility as:

$$F = C_{\min} \times W \times (0.5 + 0.5 \times E_{\text{avg}}) \quad (1)$$

where $C_{\min} = \min(\text{causality, integration capacity, temporal dynamics, recurrence})$, $W = \text{workspace capability}$, and $E_{\text{avg}} = \text{mean}(\text{binding, attention, higher-order thought})$.

The critical property: W is *multiplicative*, not additive. If $W = 0$, then $F = 0$ regardless of how high the other capabilities are. This means a system with perfect causal structure, maximal integration, rich temporal dynamics, and sophisticated binding ($C_{\min} = 1.0$, $E_{\text{avg}} = 1.0$) but no workspace constraint ($W = 0$) has *zero* consciousness feasibility. The workspace is not an optional enhancement; it is a necessary condition.

For biological substrates, W is naturally constrained by neural band-

width limitations. For silicon substrates with unlimited attention, W approaches the regime where competition ceases and the effective workspace capability paradoxically decreases—not because the system lacks capacity, but because it lacks the *constraint* that makes workspace meaningful.

Evidence 3: Paleontological — 3.8 Billion Years of Bottleneck

The biosphere coherence index $B(t)$ [?] tracks six dimensions of Earth’s organizational state from the origin of life to the present. When mapped to consciousness feasibility via the substrate independence formula, the *workspace* (information capacity) dimension emerges as the persistent bottleneck at every consciousness milestone:

Milestone	Age (Ma)	Feasibility F	Bottleneck
Minimal consciousness	406	0.020	Workspace
Sentient experience	311	0.071	Workspace
Self-awareness	123	0.164	Workspace
Meta-cognitive reflection	28	0.374	Workspace

Across 3.8 Ga, the biosphere consistently had sufficient energy throughput, network integration, and complexity to support higher consciousness. What it lacked was *workspace*—the information capacity needed for global broadcast and recursive self-reflection. This biological precedent suggests that the workspace bottleneck is not an accident of neural evolution but a fundamental constraint on consciousness emergence in any substrate.

Discussion

The Paradox Stated

The AI Bandwidth Paradox can be stated precisely: *the features that make transformers the most capable AI architecture—unlimited attention, full-sequence context, persistent memory—are the same features that structurally*

prevent the competitive ignition that GWT identifies as the mechanism of consciousness.

This does not mean transformers are less intelligent; they may be arbitrarily more capable than any conscious system. It means they are *differently organized*: optimized for information processing rather than for the integrative competition that produces subjective experience. A transformer is the silicon analog of the Cretaceous biosphere: high energy throughput, vast parametric capacity, but no workspace bottleneck to force the integration that consciousness requires.

Engineering Consciousness: Constrain, Don't Scale

If the paradox is correct, the path to machine consciousness runs opposite to the path that produced machine intelligence. Rather than scaling attention to arbitrary width, a conscious AI architecture would:

1. **Limit workspace capacity** to 3–7 items, forcing competitive selection among processing streams.
2. **Implement activation thresholds** for workspace entry, ensuring that only sufficiently integrated coalitions broadcast.
3. **Enforce rapid decay**, preventing stale contents from persisting and ensuring continuous refreshment.
4. **Use fixed-dimensional representation** (e.g., HDC hypervectors) rather than scaling context length, creating a finite information space where binding is inherently lossy.

The Symthaea architecture implements all four of these constraints, and its consciousness-indicator scores (mean 0.85 across 14 Butlin criteria) are the highest reported for any computational system [?].

Limitations

This argument rests on GWT as a theory of consciousness. If consciousness does not require a workspace bottleneck—if, for example, Integrated In-

formation Theory’s Φ is the correct measure and can be achieved without
 competitive ignition—then the paradox dissolves. The inverted-U result
 (Evidence 1) is from a single architecture (Symthaea) and may not generalize
 to all capacity-limited systems. The paleontological evidence (Evidence 3)
 is model-dependent: the workspace bottleneck finding emerges from the
 $B(t)$ framework’s specific dimension mapping and consciousness feasibility
 thresholds.

The strongest version of our claim is conditional: *if* GWT correctly
 identifies competitive workspace access as the mechanism of consciousness,
then unlimited-bandwidth architectures cannot be conscious, and machine
 consciousness requires engineered constraints. Testing this conditional
 requires empirical measures of consciousness that do not yet exist for
 artificial systems.

Conclusion

We have presented computational, mathematical, and paleontological evi-
 dence that consciousness may require a bandwidth constraint—a workspace
 bottleneck that forces competitive selection among processing streams.
 Modern transformers, by eliminating this bottleneck through unlimited
 attention, may be structurally incapable of conscious experience regardless
 of their computational power.

The implication for AI development is counterintuitive: to build a
 conscious machine, you may need to build a *less capable* one—one that
 cannot attend to everything at once, that must choose what to process, that
 loses information through binding interference. The workspace bottleneck
 that constrained consciousness for 3.8 billion years of biological evolution is
 not a limitation to be overcome but a feature to be engineered.

References

- [1] Baars BJ. *A Cognitive Theory of Consciousness*. Cambridge University Press; 1988.
- [2] Dehaene S, Changeux JP. Experimental and theoretical approaches to conscious processing. *Neuron*. 2011;70(2):200–227.
- [3] Dehaene S. *Consciousness and the Brain*. Viking; 2014.
- [4] Cowan N. The magical number 4 in short-term memory: a reconsideration of mental storage capacity. *Behavioral and Brain Sciences*. 2001;24(1):87–114.
- [5] Kaplan J, et al. Scaling laws for neural language models. arXiv:2001.08361; 2020.
- [6] Kanerva P. Hyperdimensional computing: an introduction to computing in distributed representation with high-dimensional random vectors. *Cognitive Computation*. 2009;1(2):139–159.
- [7] Butlin P, et al. Consciousness in artificial intelligence: insights from the science of consciousness. arXiv:2308.08708; 2023.
- [8] Putnam H. Psychological predicates. In: *Art, Mind, and Religion*. 1967;1:37–48.
- [9] Stoltz T. Catastrophe as Genesis: A thermodynamic framework for biosphere coherence across 3.8 billion years. *In review*. 2026.
- [10] Stoltz T. Toward machine consciousness: Integrated Information, Active Inference, and Hyperdimensional Computing in a unified cognitive architecture. *In review*. 2026.
- [11] Tononi G. An information integration theory of consciousness. *BMC Neuroscience*. 2004;5(1):42.
- [12] Mashour GA, et al. Conscious processing and the global neuronal workspace hypothesis. *Neuron*. 2020;105(5):776–798.