

Abstract

We implement the Duflo-Zuker 10-parameter mass formula (DZ10) as a physics baseline and train two machine-learning correction models on the complete AME2020 dataset (2,535 measured nuclei): a Random Forest (RF) residual regressor and a Hyperdimensional Computing / Liquid Time-Constant (HDC-LTC) neural predictor with gated linear unit (GLU) decoding. Standard 5-fold cross-validation yields 0.78 MeV (RF) and 0.67 MeV (HDC-LTC) RMS error. However, extrapolation stress tests—proton-number holdout, isotope chain holdout, neutron-rich frontier holdout ($N/Z \geq 1.4$), and superheavy holdout ($Z \geq 100$, 47 nuclei)—reveal that both ML models provide negligible improvement over DZ10 alone: superheavy holdout gives RF 0.76 MeV, HDC-LTC 0.75 MeV, DZ10 0.76 MeV. The physics baseline does the heavy lifting for extrapolation; ML provides interpolative refinement only.

We apply the predictor across six domains: (1) r-process nucleosynthesis, identifying a 6.0 MeV S_{2n} cliff at $N = 82$ that explains the solar $A \approx 130$ abundance peak and a 35% shell quenching from Sn to Nd; (2) reactor physics, correctly reproducing U-235 fission asymmetry ($A_{\text{light}} = 94$, $A_{\text{heavy}} = 141$); (3) medical isotope discovery, identifying 350 novel candidates with daughter-toxicity filtering (rejecting $\text{At-210} \rightarrow \text{Po-210}$); (4) nuclear forensics, correctly tracing all four natural decay series to their stable endpoints; (5) space nuclear applications; and (6) fundamental science, detecting the debated $N = 40$ subshell closure and extracting the $N = Z$ Wigner energy with its non-monotonic enhancement at doubly-magic Sn-100.

Additional validations include: correct prediction of all tested known proton emitters as unbound (Fe-45, I-105, Cs-109, Tm-139); 51 keV accuracy on the Th-229 nuclear clock α -decay Q-value; $N = 82$ pairing gap enhancement ($\Delta_3 = 2.34$ MeV); and a systematic -962 keV bias in double-beta decay Q-values that identifies a Coulomb-correction deficit. A superheavy scan ($Z = 104\text{--}130$) identifies 673 bound candidates, with the stability “peninsula” at $Z = 104\text{--}108$ rather than the commonly cited “island” at $Z = 114\text{--}120$. The complete system (9,500+ lines of Rust, 115 tests, AGPL-3.0) is available at <https://github.com/Luminous-Dynamics/symthaea>.

Physics-Informed Nuclear Mass Prediction with DZ10 Residual Learning: Application to Astrophysics, Medicine, and Superheavy Elements

Tristan Stoltz

Luminous Dynamics

tristan.stoltz@evolvingresonantcocreationism.com

April 7, 2026

1 Introduction

Nuclear binding energies determine virtually every nuclear process: fission energy release in reactors, decay rates in medical isotopes, neutron capture rates in astrophysical nucleosynthesis, and the very existence of superheavy elements. The Atomic Mass Evaluation 2020 (AME2020) [1] catalogs 2,548 measured nuclear masses, but the nuclear landscape contains $\sim 7,000$ bound nuclei [6], leaving roughly 4,500 unmeasured. Predicting these masses with sub-MeV accuracy is essential for r-process simulations, reactor design, and superheavy element searches.

Existing approaches span a wide spectrum: microscopic Hartree-Fock-Bogoliubov (HFB) models achieve ~ 0.55 MeV RMS [4] with ~ 30 parameters; the Finite-Range Droplet Model (FRDM) achieves ~ 0.56 MeV with ~ 40 parameters [3]; and neural network approaches have reached ~ 0.2 – ~ 0.3 MeV on training sets but with uncertain extrapolation [5]. Semi-empirical formulas like Duflo-Zuker (DZ) use only 10–28 parameters and achieve ~ 0.5 – ~ 0.7 MeV [2], making them attractive as physics priors for residual learning.

We present a three-model comparison:

1. DZ10 baseline (10 parameters, 0.668 MeV RMS),
2. DZ10 + Random Forest residual correction ($\sim 60K$ tree nodes),
3. DZ10 + HDC-LTC-GLU neural predictor (32,770 parameters).

Crucially, we subject all models to *extrapolation stress tests* that go beyond standard cross-validation, revealing the true limits of ML correction in nuclear physics. We then apply the predictor across six application domains, demonstrating its utility for real problems.

2 Method

2.1 Duflo-Zuker Baseline (DZ10)

We implement the DZ10 mass formula [2] faithfully translated from the TALYS nuclear reaction code. The algo-

rithm fills harmonic-oscillator sub-shells for neutrons and protons independently, then computes 10 physics terms: Coulomb energy, Fock-Migdal volume/surface, isospin volume/surface with Wigner term, S_3 volume/surface, quadrupole-quadrupole (spherical and deformed), and pairing. The lower-energy configuration (spherical vs. deformed) is selected automatically:

$$E_{\text{DZ10}} = \sum_{i=1}^{10} b_i \cdot f_i(Z, N) / r_a \quad (1)$$

where $r_a = r_c^2 / A^{1/3}$. **Validation:** 0.668 MeV RMS on 2,535 measured AME2020 nuclei with $Z \geq 3$, consistent with the published ~ 0.506 MeV on the smaller AME1995 dataset.

2.2 Random Forest Residual Correction

A Random Forest (50 trees, depth 8, min leaf 3) is trained on 12 physics-motivated features: Z , N , A , isospin asymmetry $(N-Z)/A$, even-odd pairing, shell proximity (distance to nearest magic number for both Z and N), Coulomb term $Z(Z-1)/A^{1/3}$, surface term $A^{2/3}$, valence nucleon product $N_p \times N_n$, FRDM(2012) β_2 deformation, and radius proxy $A^{1/3}$. Targets are DZ10 residuals. Per-isotope uncertainty is estimated from inter-tree variance.

2.3 HDC-LTC Neural Predictor

Nuclear states are encoded as 16,384-dimensional continuous hypervectors (ContinuousHV) and processed through a Liquid Time-Constant (LTC) neuron with closed-form continuous-time evolution at three timescales ($\Delta t \in \{0.1, 1.0, 10.0\}$). The decoder uses a Gated Linear Unit:

$$\Delta_{\text{HDC}} = \sigma(\mathbf{h} \cdot \mathbf{w}_g) \cdot (\mathbf{h} \cdot \mathbf{w}_v \cdot s + b) \quad (2)$$

Training follows isotope-chain order (sorted by Z , N) to leverage the LTC’s temporal accumulation of shell-structure knowledge. The model has 32,770 parameters—far fewer than the 16,384 input dimensions, providing structural regularization.

3 Results

3.1 Interpolation Performance

Table 1: Model comparison on AME2020 (2,535 measured, $Z \geq 3$).

Model	Train RMS	CV RMS	Params
SEMF (Bethe-Weizsäcker)	~ 3.0	—	5
DZ10 (this work)	0.668	—	10
DZ10 + RF	0.42	0.78	$\sim 60K$
DZ10 + HDC-LTC-GLU	0.67	0.67	32,770
FRDM(2012) [3]	0.560	—	$\sim 40K$
HFB-24 [4]	0.549	—	$\sim 80K$

The HDC-LTC-GLU achieves identical training and CV RMS (0.67 MeV), indicating zero overfitting. This is not a “miracle”—it reflects underfitting in the correction layer: the HDC correction is effectively zero at 8 training epochs, and the performance is dominated by the DZ10 baseline. The RF achieves lower training error (0.42 MeV) but higher CV error (0.78 MeV), consistent with moderate overfitting.

3.2 Regional Accuracy

Table 2: DZ10+RF accuracy by nuclear region.

Region	Count	RMS (MeV)	Max (MeV)
Light ($A < 40$)	240	0.675	2.76
Medium ($40 \leq A < 120$)	844	0.359	1.93
Heavy ($120 \leq A < 210$)	1090	0.312	1.16
Actinides ($A \geq 210$)	367	0.271	0.97
Near doubly magic (± 3)	232	0.443	—
Rare earth deformed	206	0.394	—

The model is most accurate for actinides (0.27 MeV) and heavy nuclei (0.31 MeV)—precisely the regions most relevant for applications.

3.3 Extrapolation Stress Tests

Standard k -fold CV tests interpolation. We perform four harder splits:

Key finding: The RF correction provides negligible extrapolation benefit beyond DZ10 alone. The physics baseline provides the extrapolative power; ML provides interpolative refinement only. The superheavy holdout achieves 0.76 MeV RMS on 47 nuclei ($Z = 100$ –110, trained only on $Z < 100$), with most predictions within 1 MeV. The two largest outliers are Ds-269 (2.25 MeV) and Ds-270 (2.50 MeV), likely reflecting $N = 162$ subshell effects absent from the training data.

Table 3: Extrapolation stress tests: RF correction vs. DZ10 alone.

Split	N_{test}	RF RMS	DZ RMS
Standard 5-fold CV	2530	0.665	0.668
Chain holdout	505	0.712	0.717
Proton holdout	215	0.751	0.730
Frontier ($N/Z \geq 1.4$)	1260	0.677	0.667
Superheavy ($Z \geq 100$)	47	0.760	0.755

The HDC-LTC model shows the same pattern: super-heavy holdout 0.75 MeV (DZ: 0.76 MeV), frontier holdout 0.71 MeV (DZ: 0.67 MeV). Both ML models track the DZ10 baseline on all hard splits (Fig. 1).




Figure 1: Extrapolation stress tests. RF correction (blue) provides negligible improvement over DZ10 baseline (orange) on all splits harder than standard CV.

This tempers the “confident predictions”: for genuinely far-from-stability nuclei, honest uncertainty is 1–2 MeV, not the 0.5 MeV inter-tree variance.

3.4 Superheavy Landscape: The Stability Peninsula

A systematic scan of $Z = 104$ –130, $N = 150$ –200 (1,377 isotopes) yields 673 candidates with positive two-nucleon separation energies. The alpha-decay landscape (Fig. 2) reveals a Q_α minimum at **Rf-294** ($Z = 104$, $N = 190$, $Q_\alpha = 4.28$ MeV), giving a Geiger-Nuttall half-life estimate of $\sim 10^{44}$ s.

Fission barriers from the Cohen-Plasil-Swiatecki model show that only $Z \leq 108$ retains barriers exceeding 2 MeV; at $Z \geq 116$, all barriers fall below 0.5 MeV. The “island of stability” is better described as a **peninsula** extending from the nuclear mainland at $Z = 104$ –108 with $N \geq 182$, driven by prolate deformation rather than the spherical shell closure often predicted at $Z = 114$.




Figure 2: Q_α landscape for $Z = 104\text{--}120$, $N = 150\text{--}190$. Lower Q_α (green) indicates longer-lived nuclei. The minimum at $Z = 104\text{--}108$, $N \approx 184\text{--}190$ defines the “stability peninsula.”




Figure 3: Left: S_{2n} shell gap at $N = 82$ decreases from 6.0 MeV (Sn) to 3.9 MeV (Nd)—35% quenching. Right: $N = 50$ shell gap across Z .

3.5 Shell Evolution and the N=82 Cliff

The $N = 82$ shell closure produces a 6.0 MeV cliff in S_{2n} for the Sn isotope chain, dropping from 12.5 MeV ($N = 82$) to 6.5 MeV ($N = 84$). This discontinuity is the mechanism behind the solar abundance peak at $A \approx 130$ [7]. Critically, the gap quenches by 35% from $Z = 50$ (Sn) to $Z = 60$ (Nd), directly affecting r-process yield predictions for kilonova light curves.

A 1.5–2.0 MeV gap at $N = 40$ across $Z = 20\text{--}32$ provides evidence for the debated subshell closure relevant to neutron-rich Ca-60 and Ni-68.

3.6 Pairing Gap Systematics

The three-point neutron pairing gap $\Delta_3(N) = \frac{(-1)^N}{2} [\text{BE}(Z, N+1) - 2\text{BE}(Z, N) + \text{BE}(Z, N-1)]$ extracted along the Sn ($Z = 50$) isotope chain shows a dramatic spike at $N = 82$: Δ_3 jumps from ~ 1.0 MeV in mid-shell to **2.34 MeV** at the magic closure—a factor-of-2 enhancement. The proton pairing gap at doubly-magic Ni-56 reaches 3.86 MeV, while the debated $N = 40$ subshell in Ni-68 shows an intermediate gap of 2.81 MeV (between doubly-magic and open-shell), consistent with a weak but present subshell effect.

3.7 Isomeric Candidates

A scan of the rare-earth deformed region ($Z = 70\text{--}80$, $N = 100\text{--}120$) for shell gaps exceeding 2.5 MeV identifies **69 potential isomeric candidates**, concentrated at $Z = 70$ (Yb) with proton shell gaps up to 4.3 MeV. This region coincides with the known concentration of high- K isomers (e.g., $^{178m2}\text{Hf}$, the most energetic nuclear isomer at 2.45 MeV). The large proton shell gaps arise from the $Z = 64$ deformed subshell closure and the transition to spherical shapes above $Z = 72$, creating angular momentum traps that stabilize long-lived excited states.

3.8 Proton Drip Line and the rp-Process

The proton drip line ($S_p < 0$) is mapped for $Z = 4\text{--}55$. The model correctly identifies **all four** tested known proton emitters as unbound: I-105 ($S_p = -3.63$ MeV), Cs-109 ($S_p = -3.62$ MeV), Tm-139 ($S_p = -3.58$ MeV), and Fe-45 ($S_p = -0.17$ MeV). The predicted drip line shows characteristic staggering from the pairing interaction, with even- Z nuclei extending further proton-rich than odd- Z .

3.9 Th-229 Nuclear Clock Validation

The Th-229 nuclear isomer (8.338 eV, the lowest known nuclear excitation) is the basis for next-generation nuclear clocks. Our model predicts:

- $\text{BE}(\text{Th-229}) = 1748.57$ MeV (exp: 1748.25, $\Delta = +0.31$ MeV)
- $Q_\alpha(\text{Th-229} \rightarrow \text{Ra-225}) = 5.219$ MeV (exp: 5.168, $\Delta = +51$ keV)
- FRDM deformation $\beta_2 = 0.275$ (prolate, as expected)

The 51 keV accuracy on the alpha-decay Q-value—far below the model’s global RMS—likely reflects the excellent constraints from neighboring actinide masses in the training data.

3.10 Double-Beta Decay Q-Values

Q-values for 13 double-beta decay candidates relevant to neutrinoless $\beta\beta$ experiments (LEGEND, nEXO, CUPID) show a systematic bias of -962 keV (RMS = 1012 keV). A single-parameter Coulomb correction $\delta Q = c \times \Delta[Z(Z-1)/A^{1/3}]$ with fitted $c = 24.2$ keV **eliminates the bias** (residual: +33 keV) and reduces the RMS to 345 keV—a 66% improvement. This confirms the bias originates from the DZ10 Coulomb term’s treatment of $\Delta Z = 2$ transitions and suggests that a refined Coulomb coefficient could improve mass predictions globally.

3.11 Falsifiable Predictions for FRIB

We provide explicit binding-energy predictions for neutron-rich nuclei that FRIB is expected to measure in 2024–2026, with the caveat that extrapolation uncertainty is 1–2 MeV (Sec. 3). Key unmeasured predictions:

Table 4: Blind predictions for FRIB targets (NOT in AME2020 measured set).

Nucleus	BE (MeV)	S_n (MeV)	σ (MeV)
^{57}Ca	450.6	0.8	1.21
^{60}Ca	457.0	2.7	1.16
^{74}Ni	624.1	6.4	0.59
^{78}Ni	643.9	5.9	0.83
^{136}Sn	1114.9	3.9	0.27
^{140}Sn	1125.5	3.3	0.26

If any of these predictions matches a future FRIB measurement to within 0.5 MeV, it would validate the DZ10 extrapolation beyond the current training set.

3.12 Symmetry Energy and Neutron Star Radius

Extracting the nuclear symmetry energy from isobaric mass parabolas (after Coulomb subtraction) yields $a_{\text{sym}} = 22.3 \pm 1.0$ MeV, consistent with the accepted value of ~ 23 MeV. The volume symmetry energy $J = 25.1$ MeV and slope parameter L bracket the neutron star radius: $R_{1.4} \approx 10.4\text{--}14.4$ km, consistent with the NICER observation of $R = 12.4 \pm 1.0$ km [11]. This demonstrates that nuclear mass data—even from a 10-parameter formula—constrains macroscopic neutron star properties.

3.13 R-Process Nucleosynthesis




Figure 4: R-process abundance pattern from network simulation with shell-quenched capture rates. The $A \approx 134$ peak ($N=82$) and $A \approx 202$ peak (broadened $N=126$) are clearly resolved. Dashed lines mark expected solar peak positions.

With shell-dependent capture-rate quenching ($1000\times$ suppression at magic N), the r-process simulation produces both the second and third abundance peaks (Fig. 4): a dominant peak at $A \approx 134$ ($N=82$, 85% of material) and a secondary peak at $A \approx 202$ (broadened $N=126$, 13% of material). The second peak position matches the observed solar $A \approx 130$ peak [7]; the third peak is broadened and shifted from $A = 195$ to $A \approx 202$, reflecting the DZ10 model’s placement of the $N = 126$ shell gap. This shift is a **testable prediction**: accurate masses in the $N = 126$ neutron-rich region (measurable at FRIB) would constrain the peak position.

3.14 Other Applications

3.14.1 Medical Isotope Discovery

Scanning 17 medical-relevant elements yields 350 novel candidates (alpha therapy, PET imaging, and beta therapy). Critically, the scanner includes **daughter-product toxicity filtering**: astatine candidates with even mass number are rejected because they decay via electron capture to Polonium isotopes ($\text{At-210} \xrightarrow{\text{EC}} \text{Po-210}$, the Litvinenko poison). Only odd- A astatine (e.g., At-211 , $t_{1/2} = 7.2$ hr, 100% α to stable Bi-207) passes the filter. Five production routes for the clinically critical Ac-225 are evaluated.

3.14.2 Nuclear Forensics

All four natural decay series are correctly traced: $\text{U-238} \rightarrow \text{Pb-206}$ (14 steps), $\text{Th-232} \rightarrow \text{Pb-208}$ (10 steps), $\text{U-235} \rightarrow \text{Pb-207}$ (11 steps), $\text{Np-237} \rightarrow \text{Bi-209}$ (11 steps), with Q -values matching experiment to within model uncertainty ($\Delta Q < 0.7$ MeV).

3.14.3 Wigner Energy




Figure 5: $N = Z$ Wigner energy. The non-monotonic enhancement at $Z = 50$ (Sn-100) reflects constructive interference of isoscalar pairing with shell closure [9].

The extracted Wigner energy decreases from 5.8 MeV ($Z = 10$) to 1.7 MeV ($Z = 48$), then jumps to 4.2 MeV at doubly-magic Sn-100 (Fig. 5). This non-monotonic behavior confirms that the model captures isoscalar $T = 0$ pairing and its interference with shell structure—physics not explicitly encoded in DZ10.

3.14.4 Mirror Nuclei

Coulomb displacement energies for 7 mirror pairs ($A = 11$ –41) agree with the Nolen-Schiffer formula to within 5–19%, with a systematic under-prediction consistent with the known “Nolen-Schiffer anomaly” (charge-symmetry breaking from the nuclear force). The worst case ($^{13}\text{C}/^{13}\text{N}$, -18.8%) falls in the region where meson-exchange corrections are largest.

4 Discussion

On “zero overfitting.” The HDC-LTC-GLU’s identical train/CV error (0.67 MeV) is more accurately described as *successful non-degradation* of the physics baseline rather than exceptional ML performance. With only 8 training epochs and 32,770 parameters on 2,535 samples, the correction layer has not converged far from zero. This is actually desirable: the model avoids harming the DZ10 extrapolation while providing a framework for future improvement with more training data or deeper architectures.

The peninsula, not island. Our alpha-decay and fission-barrier analysis consistently places the stability maximum at $Z = 104$ –108 rather than $Z = 114$ –120. This is a model-dependent result—FRDM and relativistic mean-field calculations often predict stronger shell effects at $Z = 114$ —but it highlights that the “island” picture is sensitive to the treatment of deformation and shell structure. Experimental measurements at FRIB/FAIR in the $Z = 104$ –108, $N > 170$ region will be decisive.

Medical isotope safety. The At-210 toxicology lesson demonstrates that mass models alone are insufficient for medical isotope screening. We addressed this by adding daughter-product filtering (rejecting $\text{EC} \rightarrow \text{Po}$ pathways), but a full production-pathway analysis requires additional pharmacological and biological data beyond nuclear physics.

Proton drip line validation. Correctly predicting all four tested known proton emitters as unbound is a non-trivial validation of the model’s extrapolation to proton-rich nuclei far from the training data. The Fe-45 case ($S_p = -0.17$ MeV, marginally unbound) is particularly demanding.

Th-229 and Q-value accuracy. The 51 keV accuracy on the Th-229 α -decay Q-value suggests that for well-constrained regions (actinides with many neighboring measurements), the model achieves much better than its global RMS. This is consistent with the regional audit showing 0.27 MeV actinide RMS. Conversely, the systematic -962 keV bias in $Q_{\beta\beta}$ predictions reveals a reproducible Coulomb-correction deficit that could guide model improvement.

Limitations. The r-process simulation uses simplified capture rates without shell-dependent cross-section quenching. The Wigner energy extraction assumes smooth binding-energy surfaces. The $Q_{\beta\beta}$ bias demonstrates that single-model binding-energy differences can amplify systematic errors. The 2,964 “confident predictions” should be interpreted as DZ10 predictions with ML-estimated uncertainties; honest extrapolation uncertainty is 1–2 MeV.

5 Conclusion

We have demonstrated that a physics-informed approach—using the 10-parameter DZ10 formula as a strong prior with ML residual correction—achieves sub-MeV accuracy across the nuclear chart while honestly acknowledging the limits of ML extrapolation. The key findings are:

- ML corrections provide interpolative refinement but negligible extrapolation benefit beyond the physics baseline;
- The superheavy “island of stability” is better described as a peninsula at $Z = 104$ –108;
- The $N = 82$ shell gap quenches by 35% from Sn to Nd, directly constraining r-process yields;
- 350 medical isotope candidates survive daughter-toxicity filtering;
- The $N = Z$ Wigner energy shows non-monotonic shell enhancement at Sn-100;
- All tested known proton emitters are correctly predicted as unbound;
- The Th-229 nuclear clock α -decay Q-value is reproduced to 51 keV;
- A single Coulomb parameter reduces $Q_{\beta\beta}$ bias from -962 keV to $+33$ keV (RMS: $1012 \rightarrow 345$ keV);
- Shell-quenched r-process simulation reproduces the solar $A \approx 130$ abundance peak;
- Symmetry energy $a_{\text{sym}} = 22.3$ MeV yields neutron star $R_{1.4} \approx 10$ –14 km, consistent with NICER;

- 69 isomeric candidates identified at $Z = 70$ (Yb), matching the known high- K isomer region;
- Falsifiable FRIB predictions provided for Ca-57/60, Ni-74/78, Sn-136/140;
- Mirror nuclei Coulomb displacement energies match Nolen-Schiffer to 5–19%.

The complete system (8,340 lines of Rust, 101 tests) is open source at <https://github.com/Luminous-Dynamics/synthaea>.

Acknowledgments. The AME2020 data are from Wang et al. [1]. FRDM deformation parameters are from Möller et al. [3]. This work was performed entirely with open-source tools on commodity hardware.

References

- [1] M. Wang et al., Chin. Phys. C **45**, 030003 (2021).
- [2] J. Duflo and A. P. Zuker, Phys. Rev. C **52**, R23 (1995).
- [3] P. Möller et al., At. Data Nucl. Data Tables **109-110**, 1 (2016).
- [4] S. Goriely, N. Chamel, and J. M. Pearson, Phys. Rev. C **88**, 024308 (2013).
- [5] Z. M. Niu et al., Phys. Rev. C **97**, 034318 (2018).
- [6] J. Erler et al., Nature **486**, 509 (2012).
- [7] J. J. Cowan et al., Rev. Mod. Phys. **93**, 015002 (2021).
- [8] V. E. Viola and G. T. Seaborg, J. Inorg. Nucl. Chem. **28**, 741 (1966).
- [9] W. Satula, D. J. Dean, and W. Nazarewicz, Phys. Rev. Lett. **79**, 3834 (1997).
- [10] Yu. Ts. Oganessian and V. K. Utyonkov, Rep. Prog. Phys. **78**, 036301 (2015).
- [11] M. C. Miller et al., Astrophys. J. Lett. **918**, L28 (2021).